

Hyperledger Fabric Peer Design

Lance Feagan

IBM China Research Lab

Topics

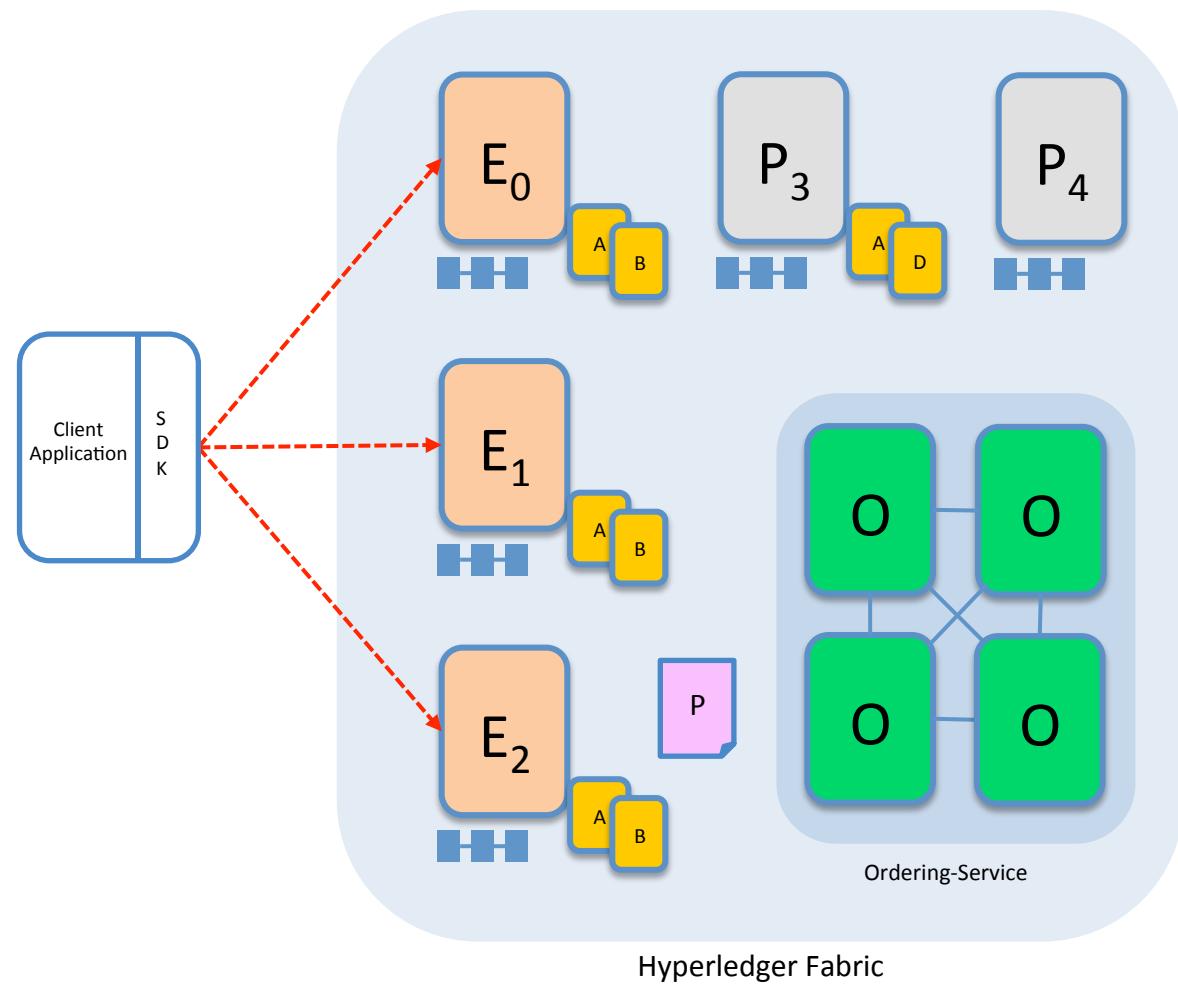
- Today
 - System Chain Code: ESCC, VSCC
 - Ledger: Design, Interfaces, Implementation
- Future
 - Gossip Protocol
 - Events

ENDORSEMENT POLICIES

Overview of Hyperledger Fabric v1 – Design Goals

- Better reflect business processes by specifying who endorses transactions
- Support broader regulatory requirements for privacy and confidentiality
- Scale the number of participants and transaction throughput
- Eliminate non deterministic transactions
- Support rich data queries of the ledger
- Dynamically upgrade the network and chaincode
- Support for multiple credential and cryptographic services for identity
- Support for "bring your own identity"

Sample transaction: Step 1/7 – Propose transaction



Application proposes transaction

Endorsement policy:

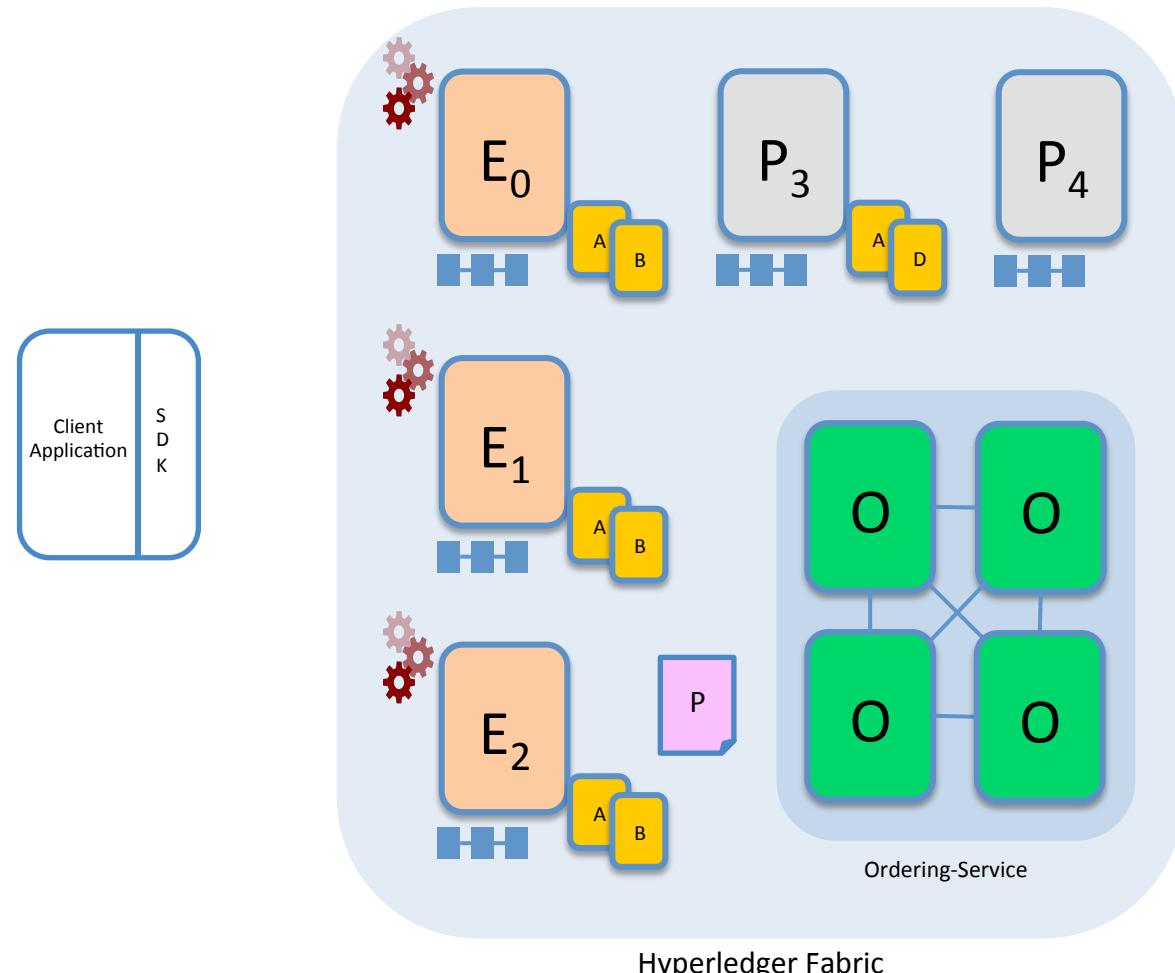
- “**E₀**, **E₁** and **E₂** must sign”
- (**P₃**, **P₄** are not part of the policy)

Client application submits a transaction proposal for **Smart Contract A**. It must target the required peers {**E₀**, **E₁**, **E₂**}

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy

Sample transaction: Step 2/7 – Execute proposal



Endorsers Execute Proposals

E_0, E_1 & E_2 will each execute the *proposed* transaction. None of these executions will update the ledger

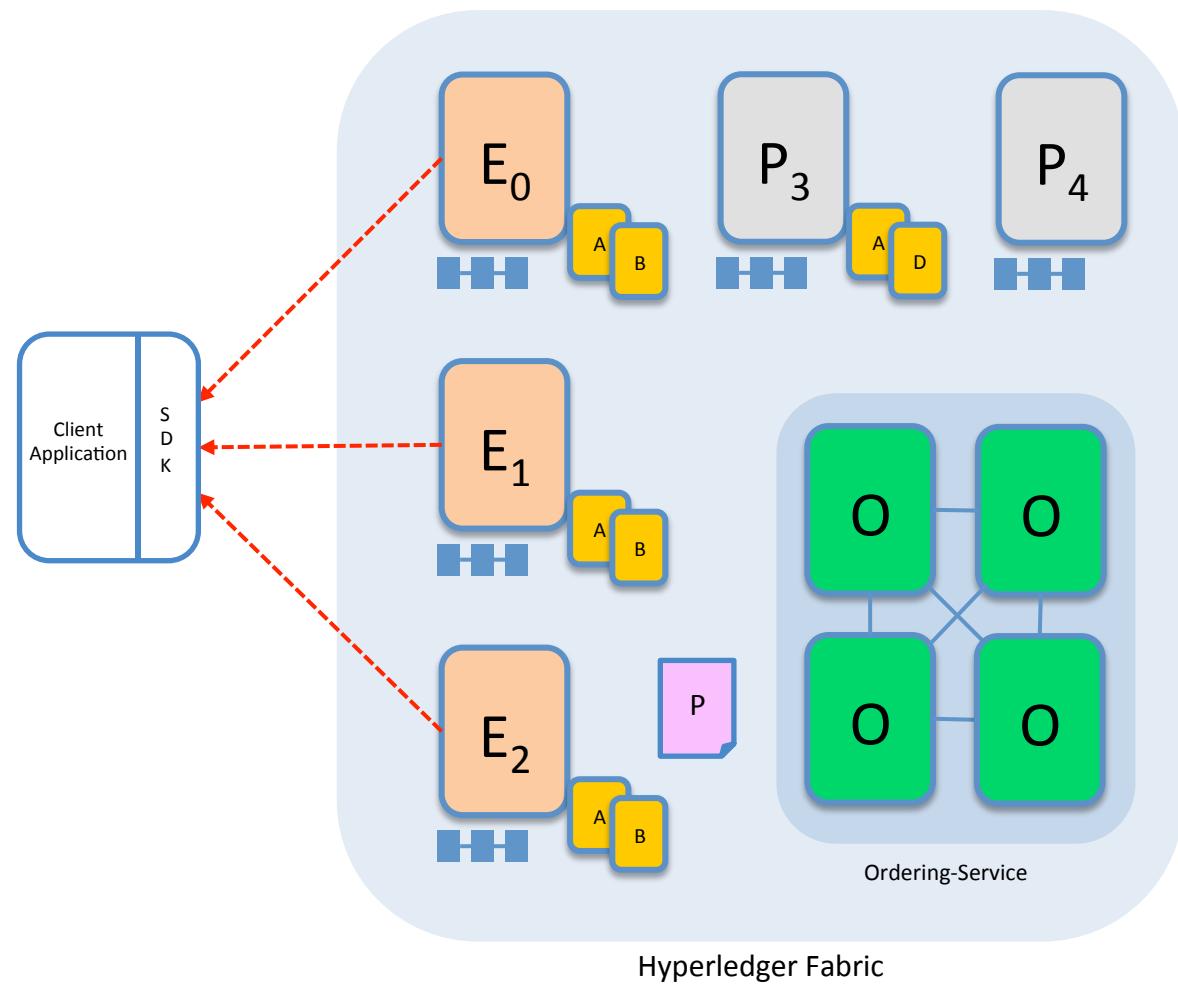
Each execution will capture the set of **Read** and **Written** data, called **RW sets**, which will now flow in the fabric.

Transactions can be signed & encrypted

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy

Sample transaction: Step 3/7 – Proposal Response



Application receives responses

RW sets are asynchronously returned to application

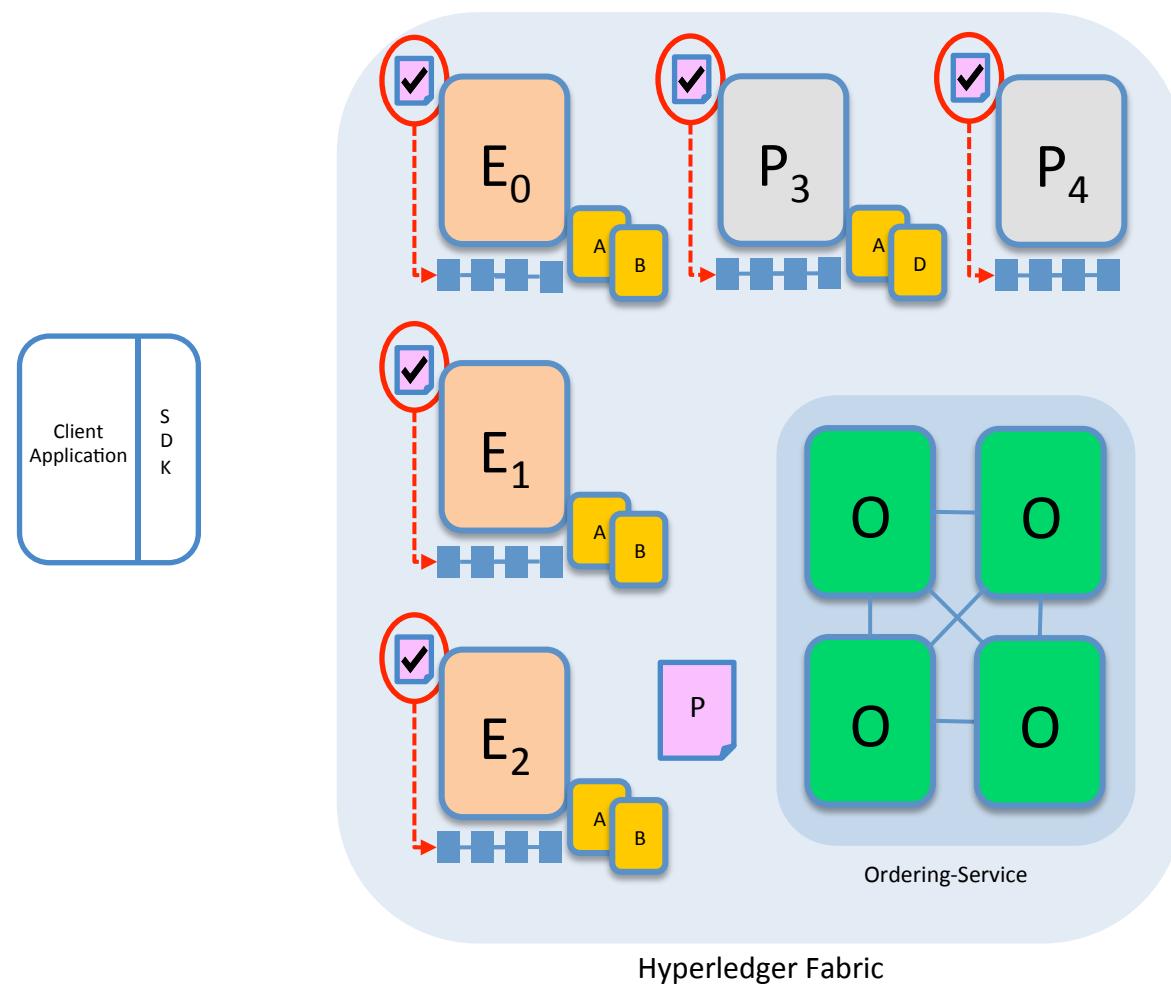
The RW sets are signed by each endorser, and also includes each record version number

(This information will be checked much later in the consensus process)

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy

Sample transaction: Step 6/7 – Validate Transaction



Committing peers validate transactions

Every committing peer validates against the endorsement policy. Also check RW sets are still valid for current world state

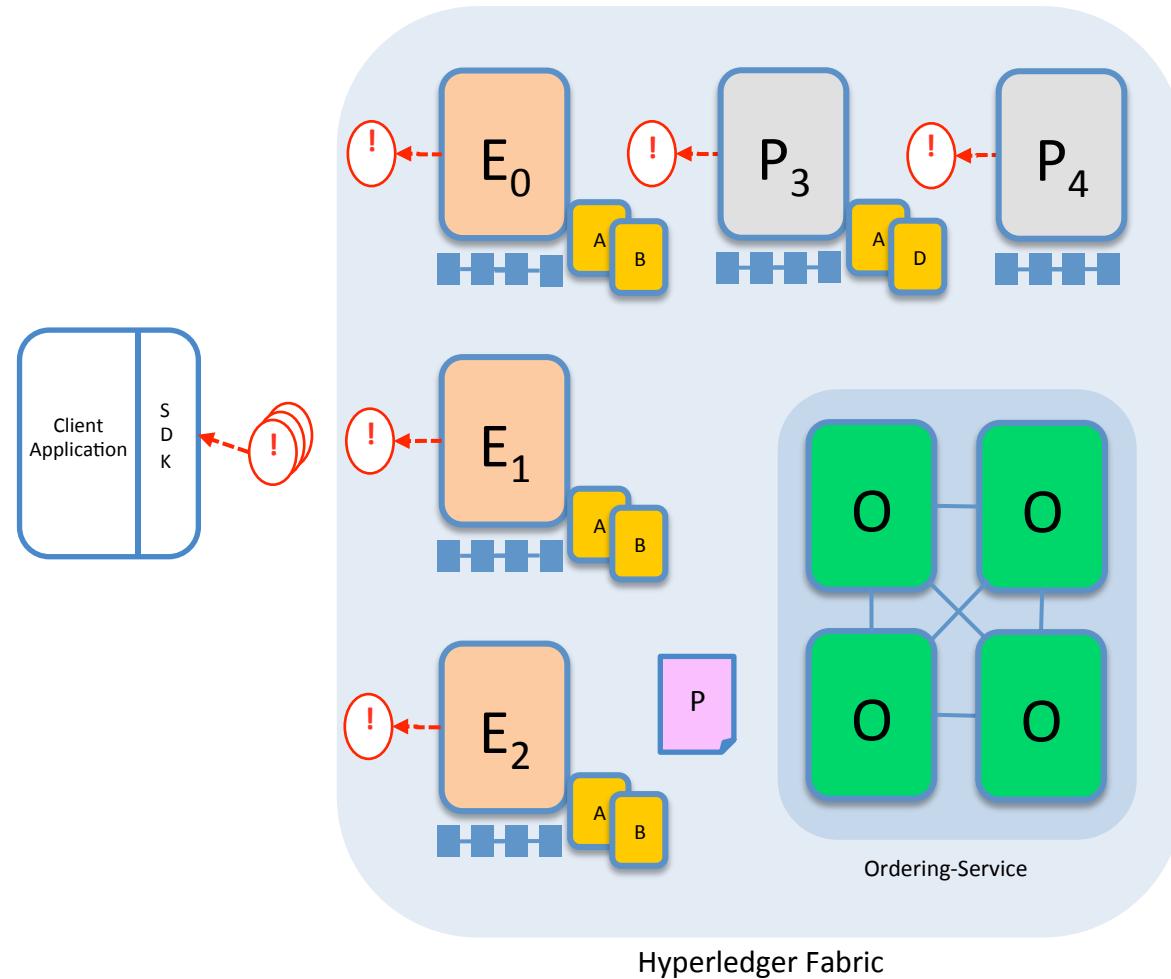
Validated transactions are applied to the world state and retained on the ledger

Invalid transactions are also retained on the ledger but do not update world state

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chaincode)		Endorsement Policy

Sample transaction: Step 7/7 – Notify Transaction



Committing peers notify applications

Applications can register to be notified when transactions succeed or fail, and when blocks are added to the ledger

Applications will be notified by each peer to which they are connected

Key:

Endorser		Ledger
Committing Peer		Application
Ordering Node		
Smart Contract (Chain code)		Endorsement Policy

Transaction Defined

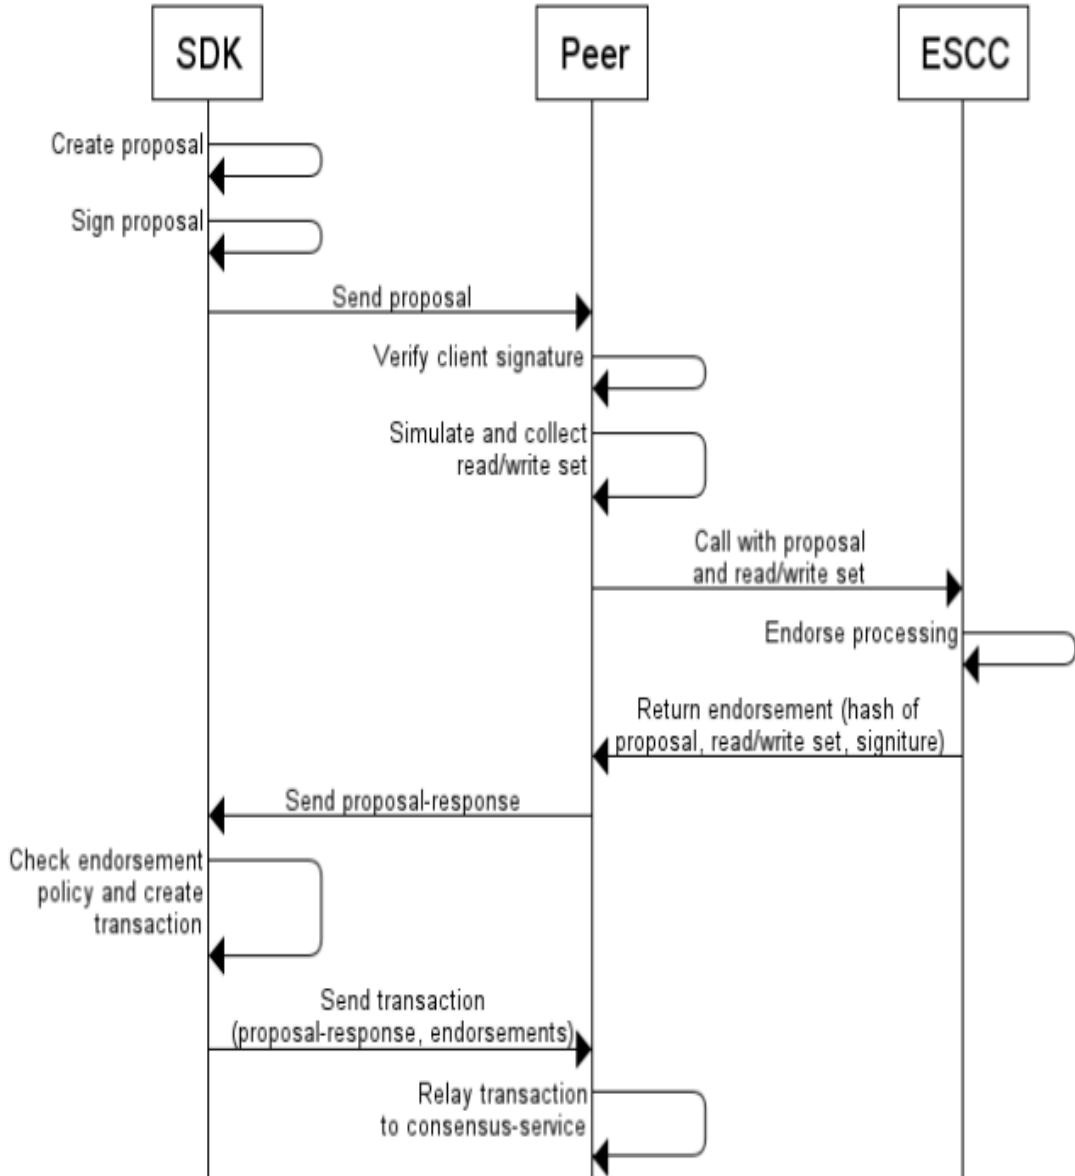
- Transaction is a chaincode function call
 - transaction : [ledger] <proposal> <endorsements>
 - proposal : [channel,] chaincode, <function-spec>
 - function-spec: function name, arguments
 - endorsements : proposal hash, simulation result, signature
- Each chaincode may be associated with an endorsement and validation system chaincode (ESCC, VSCC)
 - ESCC decides how to endorse a proposal, including simulation and application specifics
 - VSCC decides transaction validity, including correctness of endorsements

Endorsement and Validation Policy

- Each chaincode is associated with an Endorsement and Validation system chaincode, which customers may modify, add more or remove existing ones
 - A chaincode's endorsement and validation should be fixed at chaincode deployment rather than per transaction to avoid inconsistency
 - ESCC endorses a transaction, and VSCC validates the transaction according to the endorsement policy
- We should provide the following implementations of ESCC and VSCC pairs
 - At least one valid signature
 - All endorsers must sign and must be valid
 - Explicit list of signatures; the list is configured as a transaction on the ESCC
 - Percent of valid signatures (ex: 80%); the percentage is specified in a transaction
- To find ESCC and VSCC associated with a chaincode, query the chaincode deployment transaction by sending a query Proposal to the Lifecycle system chaincode

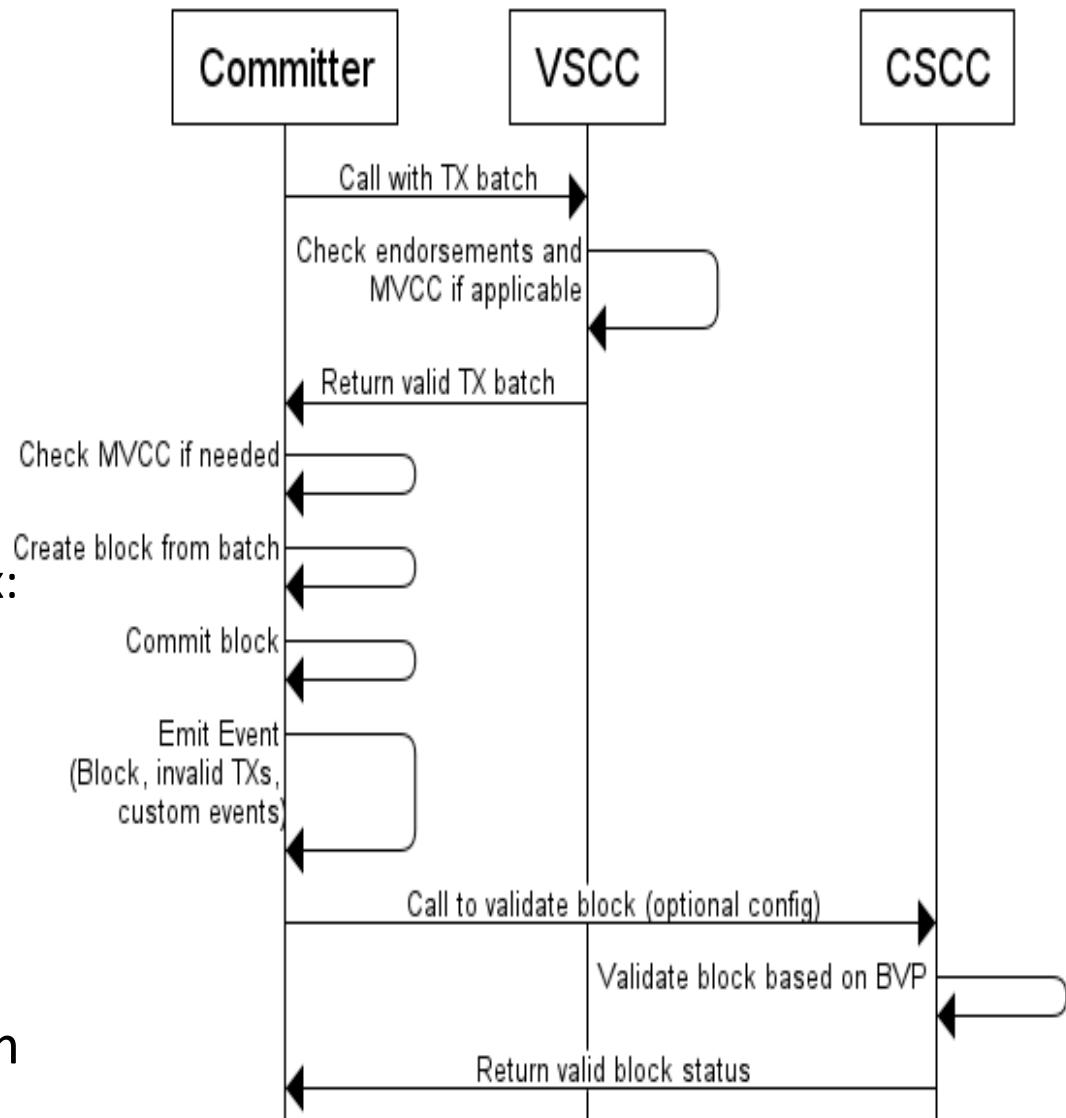
Endorse Transactions

- SDK sends Proposal to Peers based on the chaincode's endorsement policy
 - A peer may relay the proposal on client's behalf
- Endorser system chaincode (ESCC) processes the endorsement
 - ESCC provides ability to customize endorsement
 - Default logic will just sign the Proposal Response
- Client/SDK decides transaction content if endorsement satisfied



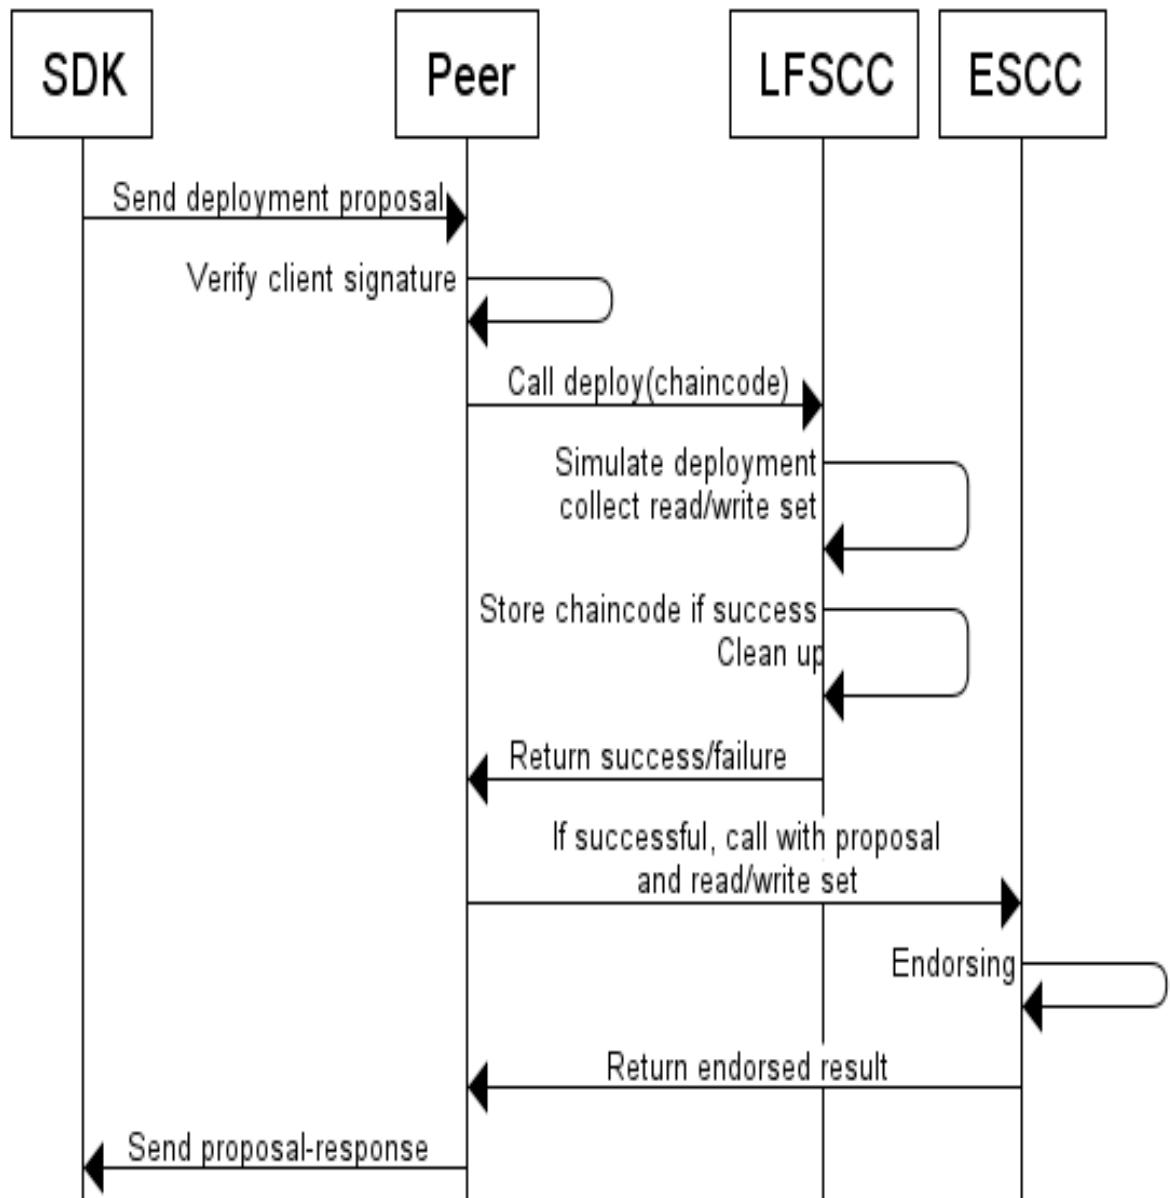
Commit Transactions

- Committing a transaction involves validating each transaction read/write-set and endorsements
- Committer calls Validator system chaincode (VSCC) to validate the batch and commit the block
 - VSCC may perform more sophisticated validation (ex: executing script OP_CHECKSIG in Bitcoin)
- Emit events (block, invalid TXs, custom)
- Block may be optionally validated (ex: checkpoint to detect faults and prune ledger) by Committer system chaincode (CSCC) based on Block Validation Policy



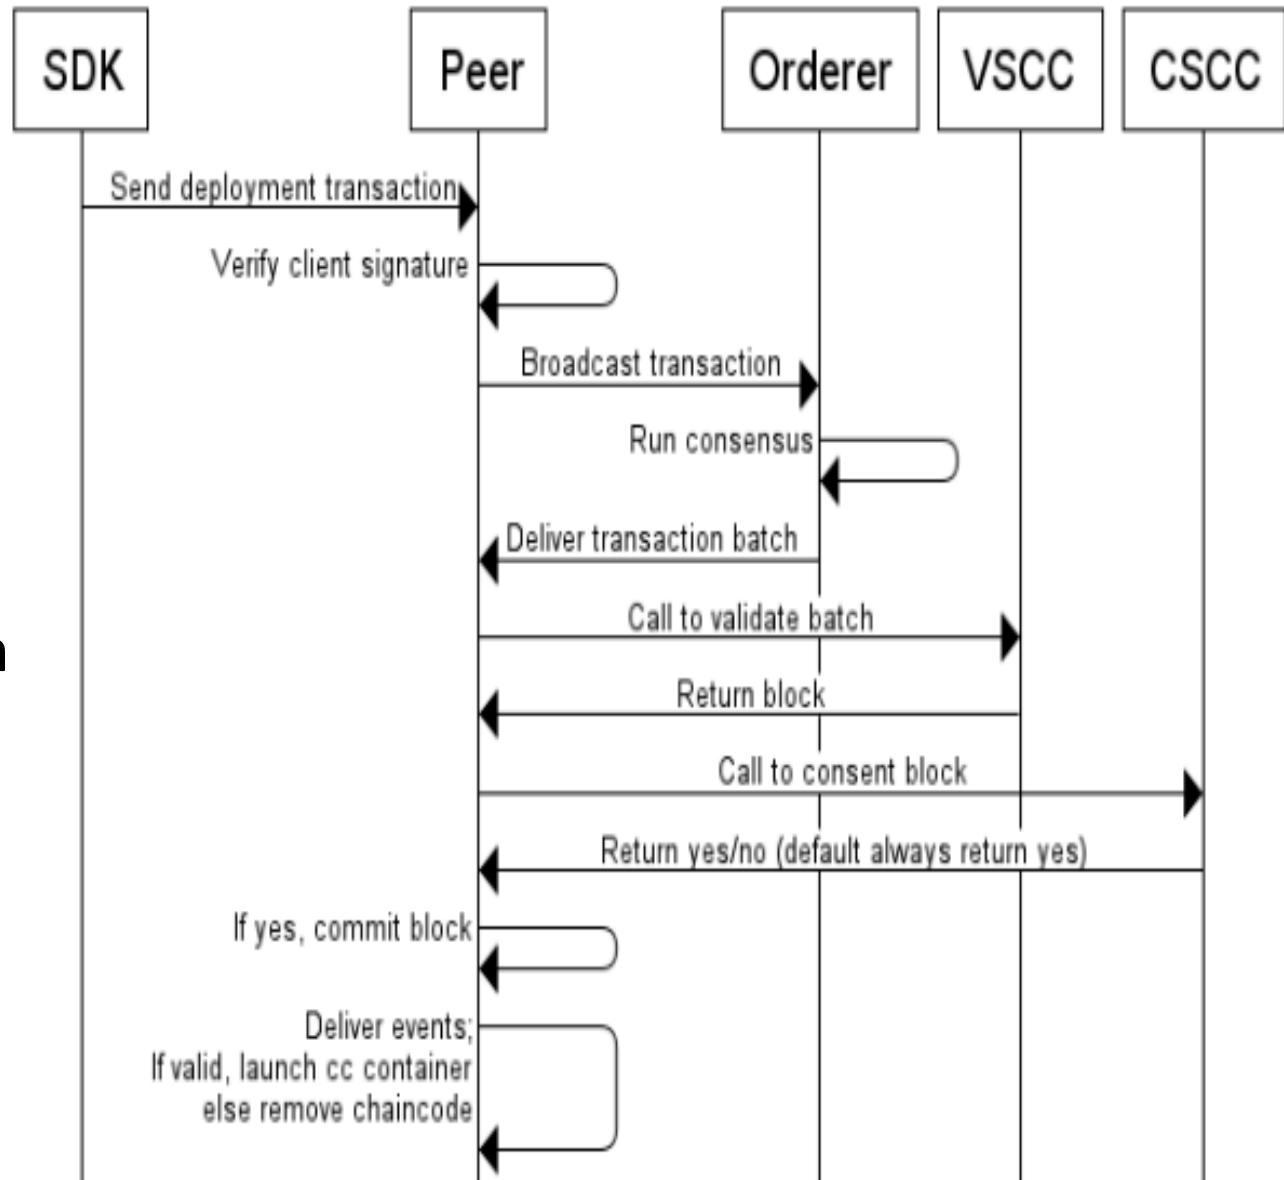
Chaincode Deployment Proposal

- Proposal is a call specification to the Life-Cycle System Chaincode LCSCC to deploy a user chaincode (UCC)
- LCSCC creates version record (version, ucc hash, name) in read-write set
- LCSCC removes container before returning



Chaincode Deployment Transaction

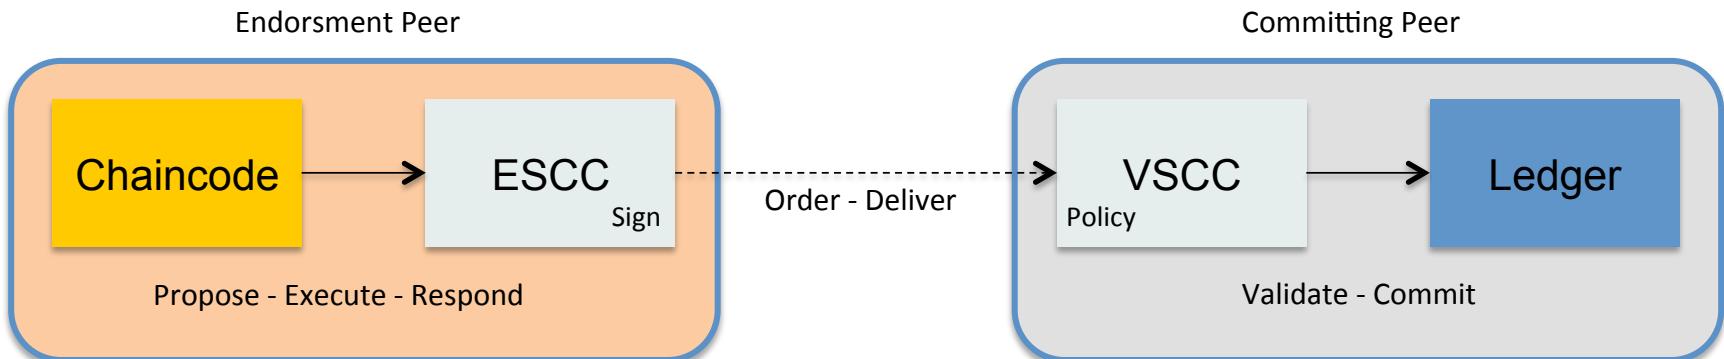
- Deployment transaction records the chaincode address and its initial values onto the ledger
- If the transaction has been successfully committed, launch the container



Endorsement Policies

An endorsement policy describes the conditions by which a transaction can be endorsed. A transaction can only be considered valid if it has been endorsed according to its policy.

- Each chaincode is associated with an Endorsement Policy
- Default implementation: Simple declarative language for the policy
- ESCC (Endorsement System ChainCode) signs the proposal response on the endorsing peer
- VSCC (Validation System ChainCode) validates the endorsements



Endorsement Policy Syntax

```
$ peer chaincode instantiate  
-C mychannel  
-n mycc  
-v 1.0  
-p chaincode_example02  
-c '{"Args":["init","a", "100",  
"b","200"]}'  
-P "AND('Org1MSP.member')"
```

This command instantiates the chaincode *mycc* on channel *mychannel* with the policy AND('Org1MSP.member')

Policy Syntax: **EXPR**(E[, E...])

Where **EXPR** is either **AND** or **OR** and **E** is either a principal or nested EXPR.

Principal Syntax: **MSP.ROLE**

Supported roles are: **member** and **admin**.

Where **MSP** is the MSP ID required, and **ROLE** is either “member” or “admin”.

Endorsement Policy Examples

Examples of policies:

- Request 1 signature from all three principals
 - AND('Org1.member', 'Org2.member', 'Org3.member')
- Request 1 signature from either one of the two principals
 - OR('Org1.member', 'Org2.member')
- Request either one signature from a member of the Org1 MSP or (1 signature from a member of the Org2 MSP and 1 signature from a member of the Org3 MSP)
 - OR('Org1.member', AND('Org2.member', 'Org3.member'))

ESCC and VSCC Source Files

- core/endorser/endorser.go:ProcessProposal
- core/committer/txvalidator/validator.go
 - Validate → validateTx →
 - VSCCValidateTx → VSCCValidateTxForCC
- common/cauthdsl/cauthdsl.go

Endorser Source 1

The screenshot shows a code editor interface with the title bar "fabric [~/gocode/src/gerrit.hyperledger.org/r/fabric] - .../core/endorser/endorser.go [fabric]". The tabs at the top include "cauthdsl.go", "policy.go", "validator.go", "endorser.go" (which is the active tab), "common.go", "chaincode.go", "instantiate.go", and "proputils.go". The code itself is a Go function named "ProcessProposal". The function takes a context.Context and a *pb.SignedProposal as parameters and returns a (*pb.ProposalResponse, error). The code handles various validation steps, including checking the proposal message, channel header, and signature header. It also checks if the proposal is for a system chaincode and returns an error if so. Finally, it checks for uniqueness of the TxID and logs the processing of the txid.

```
372
373     // ProcessProposal process the Proposal
374     func (e *Endorser) ProcessProposal(ctx context.Context, signedProp *pb.SignedProposal) (*pb.ProposalResponse, error) {
375         addr := util.ExtractRemoteAddress(ctx)
376         endorserLogger.Debug("Entering: Got request from", addr)
377         defer endorserLogger.Debugf("Exit: request from", addr)
378         // at first, we check whether the message is valid
379         prop, hdr, hdrExt, err := validation.ValidateProposalMessage(signedProp)
380         if err != nil {
381             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
382         }
383
384         chdr, err := putils.UnmarshalChannelHeader(hdr.ChannelHeader)
385         if err != nil {
386             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
387         }
388
389         shdr, err := putils.GetSignatureHeader(hdr.SignatureHeader)
390         if err != nil {
391             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
392         }
393
394         // block invocations to security-sensitive system chaincodes
395         if e.s.IsSysCCAndNotInvokableExternal(hdrExt.ChaincodeId.Name) {
396             endorserLogger.Errorf("Error: an attempt was made by %#v to invoke system chaincode %s",
397                 shdr.Creator, hdrExt.ChaincodeId.Name)
398             err = errors.Errorf("chaincode %s cannot be invoked through a proposal", hdrExt.ChaincodeId.Name)
399             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
400         }
401
402         chainID := chdr.ChannelId
403
404         // Check for uniqueness of prop.TxID with ledger
405         // Notice that ValidateProposalMessage has already verified
406         // that TxID is computed properly
407         txid := chdr.TxId
408         if txid == "" {
409             err = errors.New(message: "invalid txID. It must be different from the empty string")
410             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
411         }
412         endorserLogger.Debugf("processing txid: %s", txid)
413         if chainID != "" {
414
*Endorser.ProcessProposal(ctx context.Context, signedProp *pb.SignedProposal) (*pb.ProposalResponse, error)
```

Endorser Source 2

```
413 if chainID != "" {
414     // here we handle uniqueness check and ACLs for proposals targeting a chain
415     if _, err := e.s.GetTransactionByID(chainID, txid); err == nil {
416         return nil, errors.Errorf( format: "duplicate transaction found [%s]. Creator [%x]", txid, shdr.Creator)
417     }
418
419     // check ACL only for application chaincodes; ACLs
420     // for system chaincodes are checked elsewhere
421     if !e.s.IsSysCC(hdrExt.ChaincodeId.Name) {
422         // check that the proposal complies with the channel's writers
423         if err = e.s.CheckACL(signedProp, chdr, shdr, hdrExt); err != nil {
424             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
425         }
426     }
427 } else {
428     // chainless proposals do not/cannot affect ledger and cannot be submitted as transactions
429     // ignore uniqueness checks; also, chainless proposals are not validated using the policies
430     // of the chain since by definition there is no chain; they are validated against the local
431     // MSP of the peer instead by the call to ValidateProposalMessage above
432 }
433
434 // obtaining once the tx simulator for this proposal. This will be nil
435 // for chainless proposals
436 // Also obtain a history query executor for history queries, since tx simulator does not cover history
437 var txsim ledger.TxSimulator
438 var historyQueryExecutor ledger.HistoryQueryExecutor
439 if chainID != "" {
440     if txsim, err = e.s.GetTxSimulator(chainID, txid); err != nil {
441         return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
442     }
443     if historyQueryExecutor, err = e.s.GetHistoryQueryExecutor(chainID); err != nil {
444         return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}}, err
445     }
446     // Add the historyQueryExecutor to context
447     // TODO shouldn't we also add txsim to context here as well? Rather than passing txsim parameter
448     // around separately, since eventually it gets added to context anyways
449     ctx = context.WithValue(ctx, chaincode.HistoryQueryExecutorKey, historyQueryExecutor)
450
451     defer txsim.Done()
452 }
```

Endorser Source 3

```
455 // TODO: if the proposal has an extension, it will be of type ChaincodeAction;
456 // if it's present it means that no simulation is to be performed because
457 // we're trying to emulate a submitting peer. On the other hand, we need
458 // to validate the supplied action before endorsing it
459
460 //1 -- simulate
461 cd, res, simulationResult, ccevent, err := e.simulateProposal(ctx, chainID, txid, signedProp, prop, hdrExt.ChaincodeId, txsim)
462 if err != nil {
463     return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}, err}
464 }
465 if res != nil {
466     if res.Status >= shim.ERROR {
467         endorserLogger.Errorf( format: "simulateProposal() resulted in chaincode response status %d for txid: %s", res.Status, txid)
468         var cceventBytes []byte
469         if ccevent != nil {
470             cceventBytes, err = putils.GetBytesChaincodeEvent(ccevent)
471             if err != nil {
472                 return nil, errors.Wrap(err, message: "failed to marshal event bytes")
473             }
474         }
475         pResp, err := putils.CreateProposalResponseFailure(prop.Header, prop.Payload, res, simulationResult, cceventBytes, hdrExt.ChaincodeId, hdrExt.PayloadVisibility)
476         if err != nil {
477             return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}, err}
478         }
479
480         return pResp, &chaincodeError{ status: res.Status, msg: res.Message}
481     }
482 }
483
484 //2 -- endorse and get a marshalled ProposalResponse message
485 var pResp *pb.ProposalResponse
486
487 //TODO till we implement global ESCC, CSCC for system chaincodes
488 //chainless proposals (such as CSCC) don't have to be endorsed
489 if chainID == "" {
490     pResp = &pb.ProposalResponse{Response: res}
491 } else {
492     pResp, err = e.endorseProposal(ctx, chainID, txid, signedProp, prop, res, simulationResult, ccevent, hdrExt.PayloadVisibility, hdrExt.ChaincodeId, txsim, cd)
493     if err != nil {
494         return &pb.ProposalResponse{Response: &pb.Response{Status: 500, Message: err.Error()}, err}
495     }
496     if pResp != nil {
497         if res.Status >= shim.ERRORTHRESHOLD {
498             endorserLogger.Debugf( format: "endorseProposal() resulted in chaincode error for txid: %s", txid)
499             return pResp, &chaincodeError{ status: res.Status, msg: res.Message}
500         }
501     }
502 }
503
504 // Set the proposal response payload - it
505 // contains the "return value" from the
506 // chaincode invocation
507 pResp.Response.Payload = res.Payload
508
509 return pResp, nil
510
511 }
```

Endorser Source 4

```
283 //endorse the proposal by calling the ESCC
284 func (e *Endorser) endorseProposal(ctx context.Context, chainID string, txid string, signedProp *pb.SignedProposal,
285     proposal *pb.Proposal, response *pb.Response, simRes []byte, event *pb.ChaincodeEvent, visibility []byte,
286     ccid *pb.ChaincodeID, txsim ledger.TxSimulator, cd resourcesconfig.ChaincodeDefinition) (*pb.ProposalResponse, error) {
287     endorserLogger.Debugf(format: "Entry - txid: %s channel id: %s chaincode id: %s", txid, chainID, ccid)
288     defer endorserLogger.Debugf(format: "Exit")
289
290     isSysCC := cd == nil
291     // 1) extract the name of the escc that is requested to endorse this chaincode
292     var escc string
293     //ie, not "lscc" or system chaincodes
294     if isSysCC {
295         escc = "escc"
296     } else {
297         escc = cd.Endorsement()
298         if escc == "" { // this should never happen, LSCC always fills this field
299             panic(v: "No ESCC specified in ChaincodeData")
300         }
301     }
302
303     endorserLogger.Debugf(format: "info: escc for chaincode id %s is %s", ccid, escc)
304
305     // marshalling event bytes
306     var err error
307     var eventBytes []byte
308     if event != nil {
309         eventBytes, err = putils.GetBytesChaincodeEvent(event)
310         if err != nil {
311             return nil, errors.Wrap(err, message: "failed to marshal event bytes")
312         }
313     }
314
315     resBytes, err := putils.GetBytesResponse(response)
316     if err != nil {
317         return nil, errors.Wrap(err, message: "failed to marshal response bytes")
318     }
319
320     // set version of executing chaincode
321     if isSysCC {
322         // if we want to allow mixed fabric levels we should
323         // set syscc version to ""
324         ccid.Version = util.GetSysCCVersion()
325     } else {
326         ccid.Version = cd.CCVersion()
327     }
328
329     ccidBytes, err := putils.Marshal(ccid)
330     if err != nil {
331         return nil, errors.Wrap(err, message: "failed to marshal ChaincodeID")
332     }
```

Endorser Source 5

```
334 // 3) call the ESCC we've identified
335 // arguments:
336 // args[0] - function name (not used now)
337 // args[1] - serialized Header object
338 // args[2] - serialized ChaincodeProposalPayload object
339 // args[3] - ChaincodeID of executing chaincode
340 // args[4] - result of executing chaincode
341 // args[5] - binary blob of simulation results
342 // args[6] - serialized events
343 // args[7] - payloadVisibility
344 args := [][]byte{[]byte("")}, proposal.Header, proposal.Payload, ccidBytes, resBytes, simRes, eventBytes, visibility}
345 version := util.GetSysCCVersion()
346 ecccis := &pb.ChaincodeInvocationSpec{ChaincodeSpec: &pb.ChaincodeSpec{Type: pb.ChaincodeSpec_GOLANG, ChaincodeId: &pb.ChaincodeID{Name: escc}, Input: &pb.ChaincodeInput{Args: args}}}
347 res, _, err := e.callChaincode(ctx, chainID, version, txid, signedProp, proposal, ecccis, &pb.ChaincodeID{Name: escc}, txsim)
348 if err != nil {
349     return nil, err
350 }
351
352 if res.Status >= shim.ERRORTHRESHOLD {
353     return &pb.ProposalResponse{Response: res}, nil
354 }
355
356 prBytes := res.Payload
357 // Note that we do not extract any simulation results from
358 // the call to ESCC. This is intentional because ESCC is meant
359 // to endorse (i.e. sign) the simulation results of a chaincode,
360 // but it can't obviously sign its own. Furthermore, ESCC runs
361 // on private input (its own signing key) and so if it were to
362 // produce simulation results, they are likely to be different
363 // from other ESCCs, which would stand in the way of the
364 // endorsement process.
365
366 //3 — respond
367 pResp, err := putils.GetProposalResponse(prBytes)
368 if err != nil {
369     return nil, err
370 }
371
372 return pResp, nil
373 }
```

Validation Source 1

```
165 // Validate performs the validation of a block. The validation
166 // of each transaction in the block is performed in parallel.
167 // The approach is as follows: the committer thread starts the
168 // tx validation function in a goroutine (using a semaphore to cap
169 // the number of concurrent validating goroutines). The committer
170 // thread then reads results of validation (in order of completion
171 // of the goroutines) from the results channel. The goroutines
172 // perform the validation of the txs in the block and enqueue the
173 // validation result in the results channel. A few note-worthy facts:
174 // 1) to keep the approach simple, the committer thread enqueues
175 // all transactions in the block and then moves on to reading the
176 // results.
177 // 2) for parallel validation to work, it is important that the
178 // validation function does not change the state of the system.
179 // Otherwise the order in which validation is performed matters
180 // and we have to resort to sequential validation (or some locking).
181 // This is currently true, because the only function that affects
182 // state is when a config transaction is received, but they are
183 // guaranteed to be alone in the block. If/when this assumption
184 // is violated, this code must be changed.
185 func (v *txValidator) Validate(block *common.Block) error {
186     var err error
187     var errPos int
188
189     logger.Debug(args: "START Block Validation")
190     defer logger.Debug(args: "END Block Validation")
191     // Initialize trans as valid here, then set invalidation reason code upon validation below
192     txsfltr := ledgerUtil.NewTxValidationFlags(len(block.Data.Data))
193     // txsChaincodeNames records all the invoked chaincodes by tx in a block
194     txsChaincodeNames := make(map[int]*sysccprovider.ChaincodeInstance)
195     // upgradedChaincodes records all the chaincodes that are upgraded in a block
196     txsUpgradedChaincodes := make(map[int]*sysccprovider.ChaincodeInstance)
197     // array of txids
198     txidArray := make([]string, len(block.Data.Data))
199
200     results := make(chan *blockValidationResult)
201     go func() {
202         for tIdx, d := range block.Data.Data {
203             tIdxLcl := tIdx
204             dLcl := d
205
206             // ensure that we don't have too many concurrent validation workers
207             v.support.Acquire(context.Background(), n: 1)
208
209             go func() {
210                 defer v.support.Release(n: 1)
211
212                 validateTx(&blockValidationRequest{
213                     d: dLcl,
214                     block: block,
215                     tIdx: tIdxLcl,
216                     v: v,
217                 }, results)
218             }()
219         }()
220     }()
221
222     logger.Debugf(format: "expecting %d block validation responses", len(block.Data.Data))
223
224     // now we read responses in the order in which they come back
225     for i := 0; i < len(block.Data.Data); i++ {
226         res := <-results
227
228         if res.err != nil {
229             // if there is an error, we buffer its value, wait for
230             // all workers to complete validation and then return
231             // the error from the first tx in this block that returned an error
232             logger.Debugf(format: "got terminal error %s for idx %d", res.err, res.tIdx)
233
234             if err == nil || res.tIdx < errPos {
235                 err = res.err
236                 errPos = res.tIdx
237             }
238         } else {
239             // if there was no error, we set the txsfltr and we set the
240             // txsChaincodeNames and txsUpgradedChaincodes maps
241             logger.Debugf(format: "got result for idx %d, code %d", res.tIdx, res.validationCode)
242
243             txsfltr.SetFlag(res.tIdx, res.validationCode)
244
245             if res.validationCode == peer.TxValidationCode_VALID {
246                 if res.txsChaincodeName != nil {
247                     txsChaincodeNames[res.tIdx] = res.txsChaincodeName
248                 }
249                 if res.txsUpgradedChaincode != nil {
250                     txsUpgradedChaincodes[res.tIdx] = res.txsUpgradedChaincode
251                 }
252                 txidArray[res.tIdx] = res.txid
253             }
254         }
255     }
256
257     // if we're here, all workers have completed the validation.
258     // If there was an error we return the error from the first
259     // tx in this block that returned an error
260     if err != nil {
261         return err
262     }
263
264     // if we operate with this capability, we mark invalid any transaction that has a txid
265     // which is equal to that of a previous tx in this block
266     if v.support.Capabilities().ForbidDuplicateTXIdInBlock() {
267         markTXIdDuplicates(txidArray, txsfltr)
268     }
269
270     // if we're here, all workers have completed validation and
271     // no error was reported; we set the tx filter and return
272     // success
273
274     txsfltr = v.invalidTxsForUpgradeCC(txsChaincodeNames, txsUpgradedChaincodes, txsfltr)
275
276     // Initialize metadata structure
277     utils.InitBlockMetadata(block)
278
279     block.Metadata.Metadata[common.BlockMetadataIndex_TRANSACTIONS_FILTER] = txsfltr
280
281     return nil
282 }
```

Validation Source 2

```
302 func validateTx(req *blockValidationRequest, results chan<- *blockValidationResult) {
303     block := req.block
304     d := req.d
305     tIdx := req.tIdx
306     v := req.v
307     txID := ""
308
309     if d == nil {
310         results <- &blockValidationResult{
311             tIdx: tIdx,
312         }
313         return
314     }
315
316     if env, err := utils.GetEnvelopeFromBlock(d); err != nil {
317         logger.Warningf( format: "Error getting tx from block(%s)", err)
318         results <- &blockValidationResult{
319             tIdx: tIdx,
320             validationCode: peer.TxValidationCode_INVALID_OTHER_REASON,
321         }
322         return
323     } else if env != nil {
324         // validate the transaction: here we check that the transaction
325         // is properly formed, properly signed and that the security
326         // chain binding proposal to endorsements to tx holds. We do
327         // NOT check the validity of endorsements, though. That's a
328         // job for VSCC below
329         logger.Debugf( format: "validateTx starts for block %p env %p txn %d", block, env, tIdx)
330         defer logger.Debugf( format: "validateTx completes for block %p env %p txn %d", block, env, tIdx)
331         var payload *common.Payload
332         var err error
333         var txResult peer.TxValidationCode
334         var txsChaincodeName *sysccprovider.ChaincodeInstance
335         var txsUpgradedChaincode *sysccprovider.ChaincodeInstance
336
337         if payload, txResult = validation.ValidateTransaction(env, v.support.Capabilities()); txResult != peer.TxValidationCode_VALID {
338             logger.Errorf( format: "Invalid transaction with index %d", tIdx)
339             results <- &blockValidationResult{
340                 tIdx: tIdx,
341                 validationCode: txResult,
342             }
343             return
344         }
345
346         chdr, err := utils.UnmarshalChannelHeader(payload.Header.ChannelHeader)
347         if err != nil {
348             logger.Warningf( format: "Could not unmarshal channel header, err %s, skipping", err)
349             results <- &blockValidationResult{
350                 tIdx: tIdx,
351                 validationCode: peer.TxValidationCode_INVALID_OTHER_REASON,
352             }
353             return
354         }
```

Validation Source 3

```
356 channel := chdr.ChannelID
357 logger.Debugf( format: "Transaction is for chain %s", channel)
358
359 if !v.chainExists(channel) {
360     logger.Errorf( format: "Dropping transaction for non-existent chain %s", channel)
361     results <- &blockValidationResult{
362         tIdx:      tIdx,
363         validationCode: peer.TxValidationCode_TARGET_CHAIN_NOT_FOUND,
364     }
365     return
366 }
367
368 if common.HeaderType(chdr.Type) == common.HeaderType_ENDORSER_TRANSACTION {
369     // Check duplicate transactions
370     txID = chdr.TxId
371     if _, err := v.support.Ledger().GetTransactionByID(txID); err == nil {
372         logger.Error( args: "Duplicate transaction found, ", txID, ", skipping")
373         results <- &blockValidationResult{
374             tIdx:      tIdx,
375             validationCode: peer.TxValidationCode_DUPLICATE_TXID,
376         }
377         return
378     }
379
380     // Validate tx with vscc and policy
381     logger.Debug( args: "Validating transaction vscc tx validate")
382     err, cde := v.vscc.VSCCValidateTx(payload, d, env)
383     if err != nil {
384         logger.Errorf( format: "VSCCValidateTx for transaction txId = %s returned error %s", txID, err)
385         switch err.(type) {
386             case *VSCCExecutionFailureError:
387                 results <- &blockValidationResult{
388                     tIdx: tIdx,
389                     err:  err,
390                 }
391                 return
392             case *VSCCInfoLookupFailureError:
393                 results <- &blockValidationResult{
394                     tIdx: tIdx,
395                     err:  err,
396                 }
397                 return
398             default:
399                 results <- &blockValidationResult{
400                     tIdx:      tIdx,
401                     validationCode: cde,
402                 }
403                 return
404             }
405         }
406
407     invokeCC, upgradeCC, err := v.getTxCInstance(payload)
408     if err != nil {
409         logger.Errorf( format: "Get chaincode instance from transaction txId = %s returned error %s", txID, err)
410         results <- &blockValidationResult{
411             tIdx:      tIdx,
412             validationCode: peer.TxValidationCode_INVALID_OTHER_REASON,
413         }
414     }
415 }
```

Validation Source 4

```
416 txsChaincodeName = invokeCC
417 if upgradeCC != nil {
418     logger.Infof( format: "Find chaincode upgrade transaction for chaincode %s on chain %s with new version %s",
419     txsUpgradedChaincode = upgradeCC
420 }
421 } else if common.HeaderType(chdr.Type) == common.HeaderType_CONFIG {
422     configEnvelope, err := configtx.UnmarshalConfigEnvelope(payload.Data)
423     if err != nil {
424         err := fmt.Errorf( format: "Error unmarshaling config which passed initial validity checks: %s", err)
425         logger.Critical(err)
426         results <- &blockValidationResult{
427             tIdx: tIdx,
428             err: err,
429         }
430         return
431     }
432
433     if err := v.support.Apply(configEnvelope); err != nil {
434         err := fmt.Errorf( format: "Error validating config which passed initial validity checks: %s", err)
435         logger.Critical(err)
436         results <- &blockValidationResult{
437             tIdx: tIdx,
438             err: err,
439         }
440         return
441     }
442     logger.Debugf( format: "config transaction received for chain %s", channel)
443 } else if common.HeaderType(chdr.Type) == common.HeaderType_PEER_RESOURCE_UPDATE {
444     // FIXME: in the context of FAB-7341, we should introduce validation
445     //        for this kind of transaction here. For now we just ignore this
446     //        type of transaction and delegate its validation to other components
447
448     results <- &blockValidationResult{
449         tIdx: tIdx,
450         err: nil,
451     }
452     return
453 } else {
454     logger.Warningf( format: "Unknown transaction type [%s] in block number [%d] transaction index [%d]",
455     common.HeaderType(chdr.Type), block.Header.Number, tIdx)
456     results <- &blockValidationResult{
457         tIdx: tIdx,
458         validationCode: peer.TxValidationCode_UNKNOWN_TX_TYPE,
459     }
460     return
461 }
462
463 if _, err := proto.Marshal(env); err != nil {
464     logger.Warningf( format: "Cannot marshal transaction due to %s", err)
465     results <- &blockValidationResult{
466         tIdx: tIdx,
467         validationCode: peer.TxValidationCode_MARSHAL_TX_ERROR,
468     }
469     return
470 }
```

Validation Source 5

```
471 // Succeeded to pass down here, transaction is valid
472 results <- &blockValidationResult{
473     tIdx:                      tIdx,
474     txsChaincodeName:          txsChaincodeName,
475     txsUpgradedChaincode:      txsUpgradedChaincode,
476     validationCode:            peer.TxValidationCode_VALID,
477     txid:                      txID,
478 }
479     return
480 } else {
481     logger.Warning( args: "Nil tx from block")
482     results <- &blockValidationResult{
483         tIdx:                      tIdx,
484         validationCode:             peer.TxValidationCode_NIL_ENVELOPE,
485     }
486     return
487 }
488 }
```

Validation Source 6

```
680 ① func (v *vsccValidatorImpl) VSCCValidateTx(payload *common.Payload, envBytes []byte, env *common.Envelope) (error, peer.TxValidationCode) {
681    logger.Debugf( format: "VSCCValidateTx starts for env %p envbytes %p", env, envBytes)
682    defer logger.Debugf( format: "VSCCValidateTx completes for env %p envbytes %p", env, envBytes)
683
684    // get header extensions so we have the chaincode ID
685    hdrExt, err := utils.GetChaincodeHeaderExtension(payload.Header)
686    if err != nil {
687        return err, peer.TxValidationCode_BAD_HEADER_EXTENSION
688    }
689
690    // get channel header
691    chdr, err := utils.UnmarshalChannelHeader(payload.Header.ChannelHeader)
692    if err != nil {
693        return err, peer.TxValidationCode_BAD_CHANNEL_HEADER
694    }
695
696    /* obtain the list of namespaces we're writing stuff to;
697     at first, we establish a few facts about this invocation:
698     1) which namespaces does it write to?
699     2) does it write to LSCC's namespace?
700     3) does it write to any cc that cannot be invoked? */
701    wrNamespace := []string{}
702    writesToLSCC := false
703    writesToNonInvokableSCC := false
704    respPayload, err := utils.GetActionFromEnvelope(envBytes)
705    if err != nil {
706        return fmt.Errorf( format: "GetActionFromEnvelope failed, error %s", err), peer.TxValidationCode_BAD_RESPONSE_PAYLOAD
707    }
708    txRwSet := &rwsetutil.TxRwSet{}
709    if err = txRwSet.FromProtoBytes(respPayload.Results); err != nil {
710        return fmt.Errorf( format: "txRwSet.FromProtoBytes failed, error %s", err), peer.TxValidationCode_BAD_RWSET
711    }
712    for _, ns := range txRwSet.NsRwSets {
713        if v.txWritesToNamespace(ns) {
714            wrNamespace = append(wrNamespace, ns.NameSpace)
715
716            if !writesToLSCC && ns.NameSpace == "lsc" {
717                writesToLSCC = true
718            }
719
720            if !writesToNonInvokableSCC && v.sccprovider.IsSysCCAndNotInvokableCC2CC(ns.NameSpace) {
721                writesToNonInvokableSCC = true
722            }
723
724            if !writesToNonInvokableSCC && v.sccprovider.IsSysCCAndNotInvokableExternal(ns.NameSpace) {
725                writesToNonInvokableSCC = true
726            }
727        }
728    }
}
```

Validation Source 7

```
730 // get name and version of the cc we invoked
731 ccID := hdrExt.ChaincodeId.Name
732 ccVer := respPayload.ChaincodeId.Version
733
734 // sanity check on ccID
735 if ccID == "" {
736     err := fmt.Errorf("invalid chaincode ID")
737     logger.Errorf("format: %s", err)
738     return err, peer.TxValidationCode_INVALID_OTHER_REASON
739 }
740 if ccID != respPayload.ChaincodeId.Name {
741     err := fmt.Errorf("inconsistent ccid info (%s/%s)", ccID, respPayload.ChaincodeId.Name)
742     logger.Errorf("format: %s", err)
743     return err, peer.TxValidationCode_INVALID_OTHER_REASON
744 }
745 // sanity check on ccver
746 if ccVer == "" {
747     err := fmt.Errorf("invalid chaincode version")
748     logger.Errorf("format: %s", err)
749     return err, peer.TxValidationCode_INVALID_OTHER_REASON
750 }
751
752 // we've gathered all the info required to proceed to validation;
753 // validation will behave differently depending on the type of
754 // chaincode (system vs. application)
755
756 if !v.sccprovider.IsSysCC(ccID) {
757     // if we're here, we know this is an invocation of an application chaincode;
758     // first of all, we make sure that:
759     // 1) we don't write to LSCC - an application chaincode is free to invoke LSCC
760     //    for instance to get information about itself or another chaincode; however
761     //    these legitimate invocations only ready from LSCC's namespace; currently
762     //    only two functions of LSCC write to its namespace: deploy and upgrade and
763     //    neither should be used by an application chaincode
764     if writesToLSCC {
765         return fmt.Errorf("Chaincode %s attempted to write to the namespace of LSCC", ccID),
766             peer.TxValidationCode_ILLEGAL_WRITESET
767     }
768     // 2) we don't write to the namespace of a chaincode that we cannot invoke - if
769     //    the chaincode cannot be invoked in the first place, there's no legitimate
770     //    way in which a transaction has a write set that writes to it; additionally
771     //    we don't have any means of verifying whether the transaction had the rights
772     //    to perform that write operation because in v1, system chaincodes do not have
773     //    any endorsement policies to speak of. So if the chaincode can't be invoked
774     //    it can't be written to by an invocation of an application chaincode
775     if writesToNonInvokableSCC {
776         return fmt.Errorf("Chaincode %s attempted to write to the namespace of a system chaincode that cannot be invoked", ccID),
777             peer.TxValidationCode_ILLEGAL_WRITESET
778     }
779
780 // validate *EACH* read write set according to its chaincode's endorsement policy
781 for _, ns := range wrNamespace {
782     // Get latest chaincode version, vscc and validate policy
783     txcc, vscc, policy, err := v.GetInfoForValidate(chdr.TxId, chdr.ChannelId, ns)
784     if err != nil {
785         logger.Errorf("GetInfoForValidate for txId = %s returned error %s", chdr.TxId, err)
786         return err, peer.TxValidationCode_INVALID_OTHER_REASON
787     }
788 }
```

Validation Source 8

```
789 // if the namespace corresponds to the cc that was originally
790 // invoked, we check that the version of the cc that was
791 // invoked corresponds to the version that lscc has returned
792 if ns == ccID && txcc.ChaincodeVersion != ccVer {
793     err := fmt.Errorf( format: "Chaincode %s:%s didn't match %s:%s in lscc", ccID, ccVer, chdr.ChannelId, txcc.ChaincodeName, txcc.ChaincodeVersion, chdr.ChannelId)
794     logger.Errorf(err.Error())
795     return err, peer.TxValidationCode_EXPIRED_CHAINCODE
796 }
797
798 // do VSCC validation
799 if err = v.VSCCValidateTxForCC(envBytes, chdr.TxId, chdr.ChannelId, vscc.ChaincodeName, vscc.ChaincodeVersion, policy); err != nil {
800     switch err.(type) {
801     case *VSCCEndorsementPolicyError:
802         return err, peer.TxValidationCode_ENDORSEMENT_POLICY_FAILURE
803     default:
804         return err, peer.TxValidationCode_INVALID_OTHER_REASON
805     }
806 }
807
808 } else {
809     // make sure that we can invoke this system chaincode - if the chaincode
810     // cannot be invoked through a proposal to this peer, we have to drop the
811     // transaction; if we didn't, we wouldn't know how to decide whether it's
812     // valid or not because in v1, system chaincodes have no endorsement policy
813     if v.sccprovider.IsSysCCAndNotInvokableExternal(ccID) {
814         return fmt.Errorf( format: "Committing an invocation of cc %s is illegal", ccID),
815         peer.TxValidationCode_ILLEGAL_WRITESET
816     }
817
818     // Get latest chaincode version, vscc and validate policy
819     _, vscc, policy, err := v.GetInfoForValidate(chdr.TxId, chdr.ChannelId, ccID)
820     if err != nil {
821         logger.Errorf( format: "GetInfoForValidate for txId = %s returned error %s", chdr.TxId, err)
822         return err, peer.TxValidationCode_INVALID_OTHER_REASON
823     }
824
825     // validate the transaction as an invocation of this system chaincode;
826     // vscc will have to do custom validation for this system chaincode
827     // currently, VSCC does custom validation for LSCC only; if an hlf
828     // user creates a new system chaincode which is invokable from the outside
829     // they have to modify VSCC to provide appropriate validation
830     if err = v.VSCCValidateTxForCC(envBytes, chdr.TxId, vscc.ChainID, vscc.ChaincodeName, vscc.ChaincodeVersion, policy); err != nil {
831         switch err.(type) {
832         case *VSCCEndorsementPolicyError:
833             return err, peer.TxValidationCode_ENDORSEMENT_POLICY_FAILURE
834         default:
835             return err, peer.TxValidationCode_INVALID_OTHER_REASON
836         }
837     }
838
839     return nil, peer.TxValidationCode_VALID
840
841 }
```

Validation Source 9

```
843 func (v *vsccValidatorImpl) VSCCValidateTxForCC(envBytes []byte, txid, chid, vsccName, vsccVer string, policy []byte) error {
844     logger.Debugf(format: "VSCCValidateTxForCC starts for envbytes %p", envBytes)
845     defer logger.Debugf(format: "VSCCValidateTxForCC completes for envbytes %p", envBytes)
846     ctxt, txsim, err := v.ccprovider.GetContext(v.support.Ledger(), txid)
847     if err != nil {
848         msg := fmt.Sprintf(format: "Cannot obtain context for txid=%s, err %s", txid, err)
849         logger.Errorf(msg)
850         return &VSCCExecutionFailureError{reason: msg}
851     }
852     defer txsim.Done()
853
854     // build arguments for VSCC invocation
855     // args[0] - function name (not used now)
856     // args[1] - serialized Envelope
857     // args[2] - serialized policy
858     args := [][]byte{[]byte(""), envBytes, policy}
859
860     // get context to invoke VSCC
861     vscctxid := coreUtil.GenerateUUID()
862     cccid := v.ccprovider.GetCCContext(chid, vsccName, vsccVer, vscctxid, syscc: true, signedProp: nil, prop: nil)
863
864     // invoke VSCC
865     logger.Debug(args: "Invoking VSCC txid", txid, "chainID", chid)
866     res, _, err := v.ccprovider.ExecuteChaincode(ctxt, cccid, args)
867     if err != nil {
868         msg := fmt.Sprintf(format: "Invoke VSCC failed for transaction txid=%s, error %s", txid, err)
869         logger.Errorf(msg)
870         return &VSCCExecutionFailureError{reason: msg}
871     }
872     if res.Status != shim.OK {
873         logger.Errorf(format: "VSCC check failed for transaction txid=%s, error %s", txid, res.Message)
874         return &VSCCEndorsementPolicyError{reason: fmt.Sprintf(format: "%s", res.Message)}
875     }
876
877     return nil
878 }
```

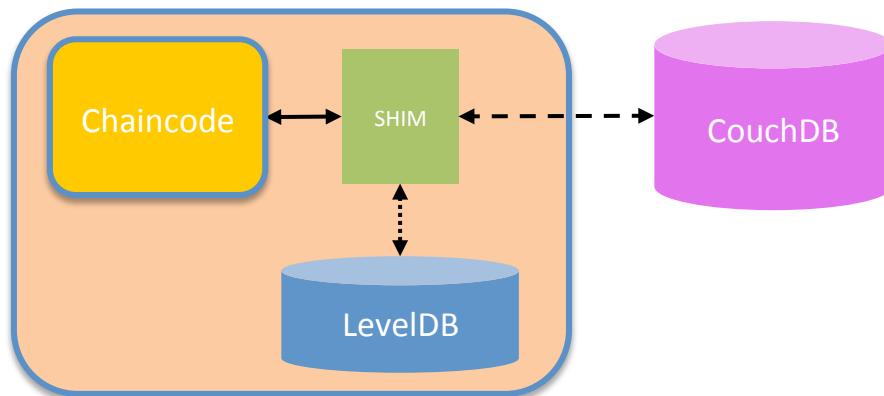
Pluggable, Queryable World State

Overview of Hyperledger Fabric v1 – Design Goals

- Better reflect business processes by specifying who endorses transactions
- Support broader regulatory requirements for privacy and confidentiality
- Scale the number of participants and transaction throughput
- Eliminate non deterministic transactions
- **Support rich data queries of the ledger**
- Dynamically upgrade the network and chaincode
- Support for multiple credential and cryptographic services for identity
- Support for "bring your own identity"

World State Database

- Pluggable worldstate database
- Default embedded key/value implementation using LevelDB
 - Support for keyed queries, but cannot query on value
- Support for Apache CouchDB
 - Full query support on key and value (JSON documents)
 - Meets a large range of chaincode, auditing, and reporting requirements
 - Will support reporting and analytics via data replication to an analytics engine such as Spark (future)
 - Id/document data model compatible with existing chaincode key/value programming model



Ledger v1 Objectives

Support v1 endorsement/consensus model - separation of simulation (chaincode execution) and block validation/commit

- Chaincode execution simulated on ‘endorsing’ peers (subset of ‘committing’ peers)
- Transaction validated and committed on all ‘committing’ peers
- Parallel simulation enabled on endorsers for improved concurrency and scalability

Persist transaction read/write sets on the blockchain

- Immutability, Auditing, Provenance

Remove dependence on RocksDB

- Utilize LevelDB instead

Optimize data storage for blockchain use patterns

- New file-based blockchain ledger for immutable transaction log
- LevelDB indexes against file-based ledger for efficient lookups
- LevelDB key/value state database for transaction execution (by default)

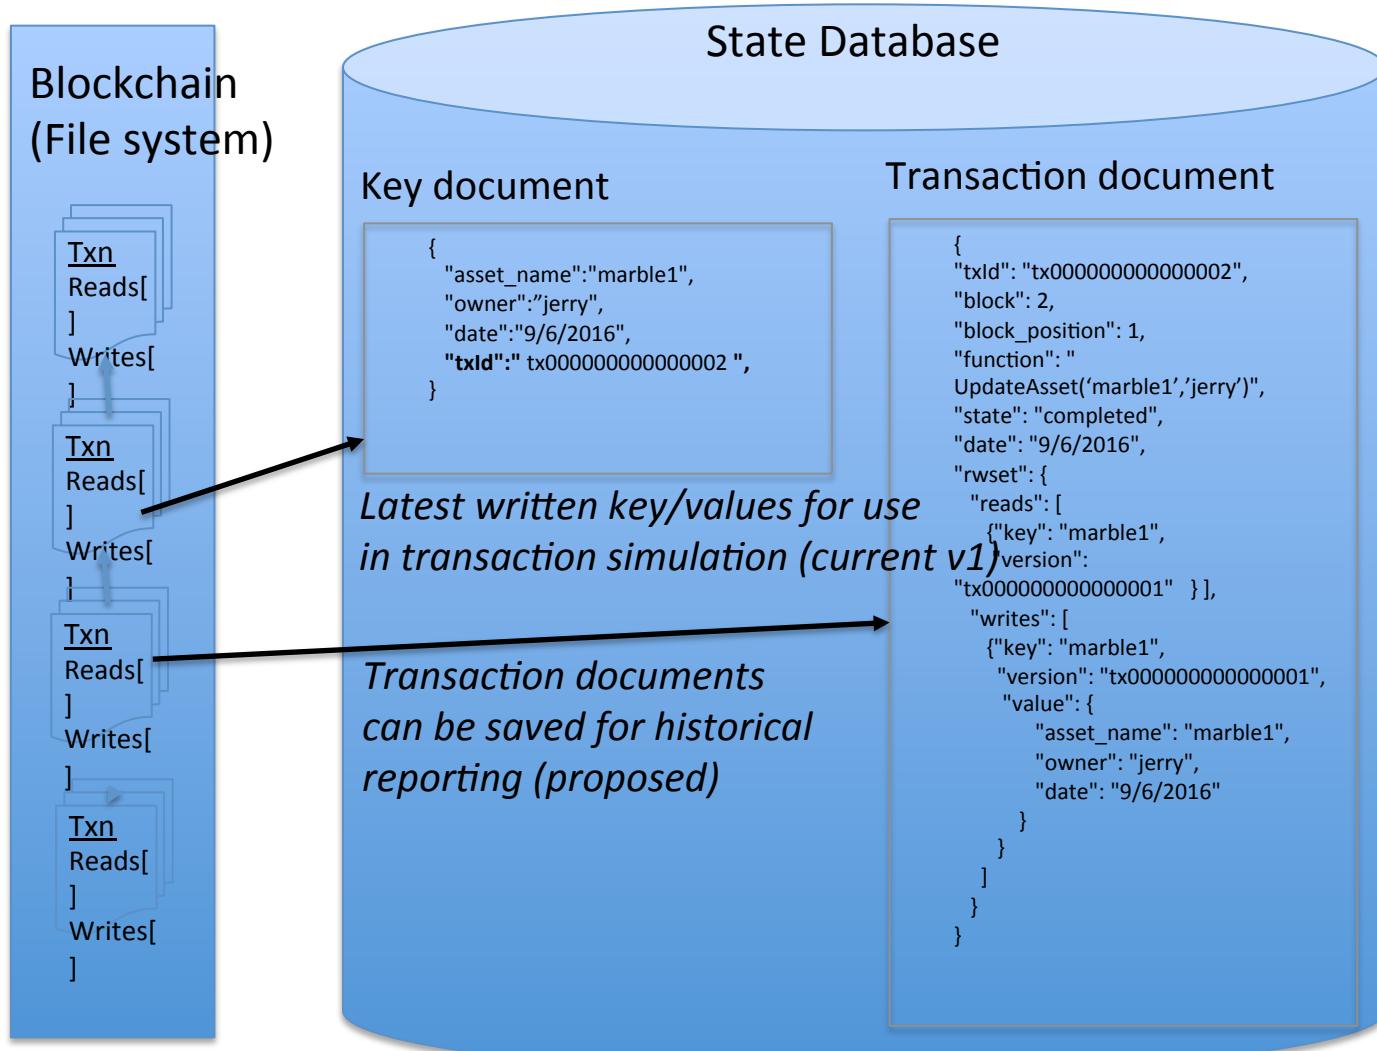
Enrich query capability of data in the blockchain

- Efficient non-key queries
- Historical queries (simple provenance scenarios)

Support for plugging in external state databases

- First external database is CouchDB – supports rich data query when modeling chaincode data as JSON

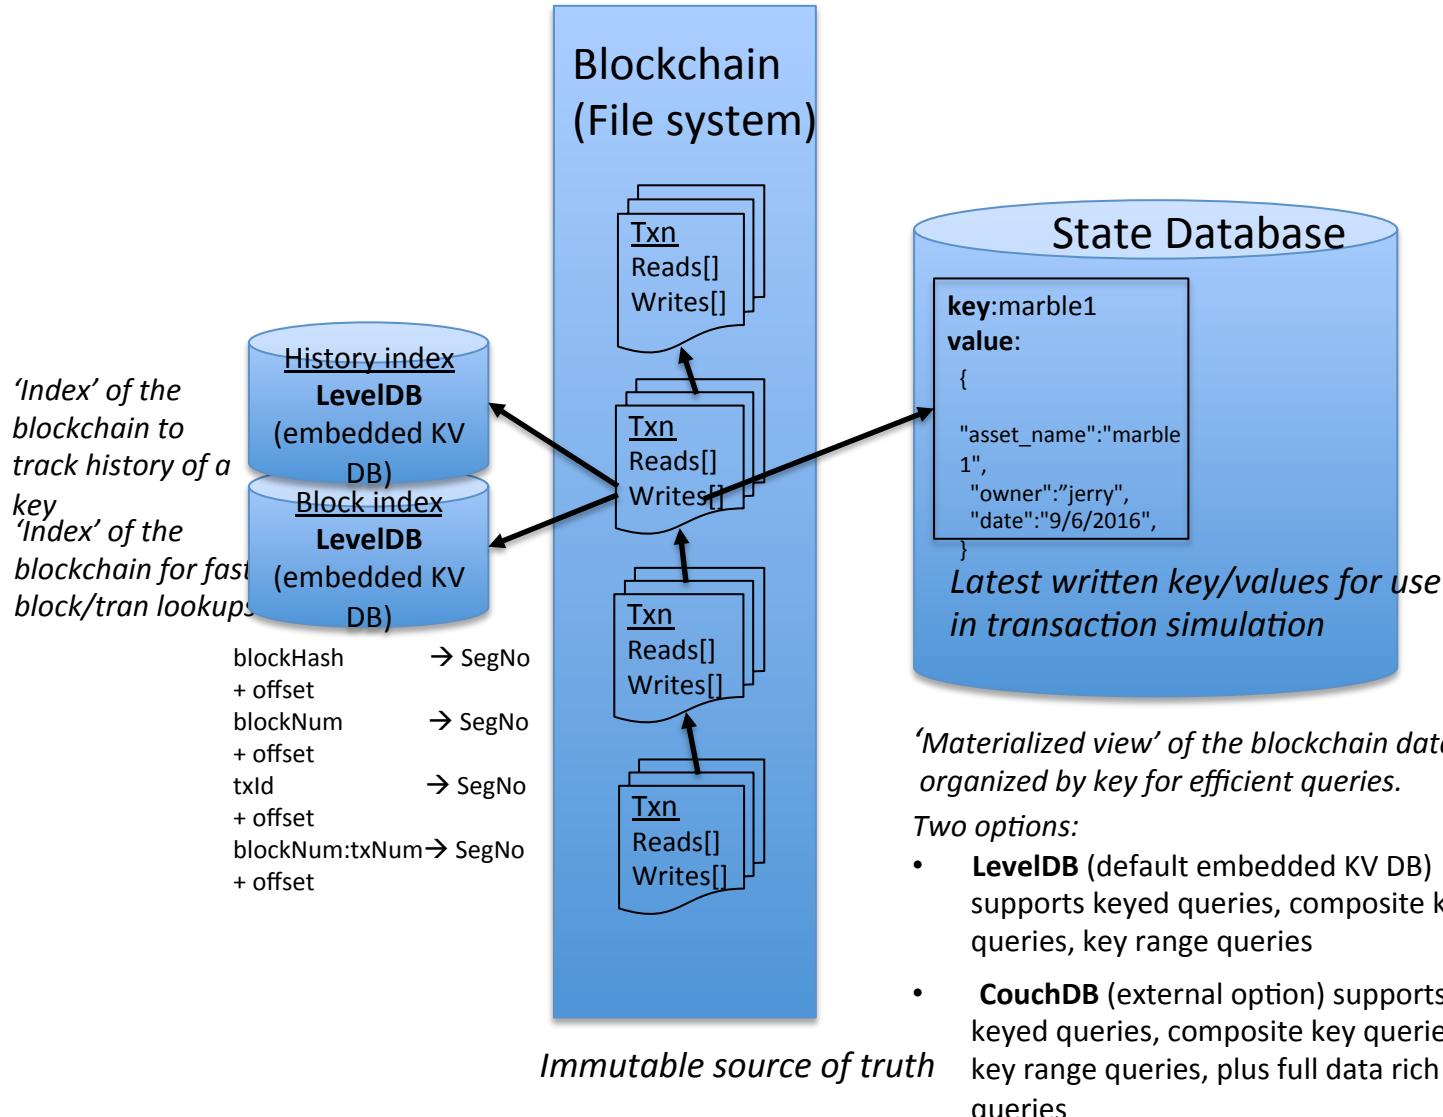
Ledger v1



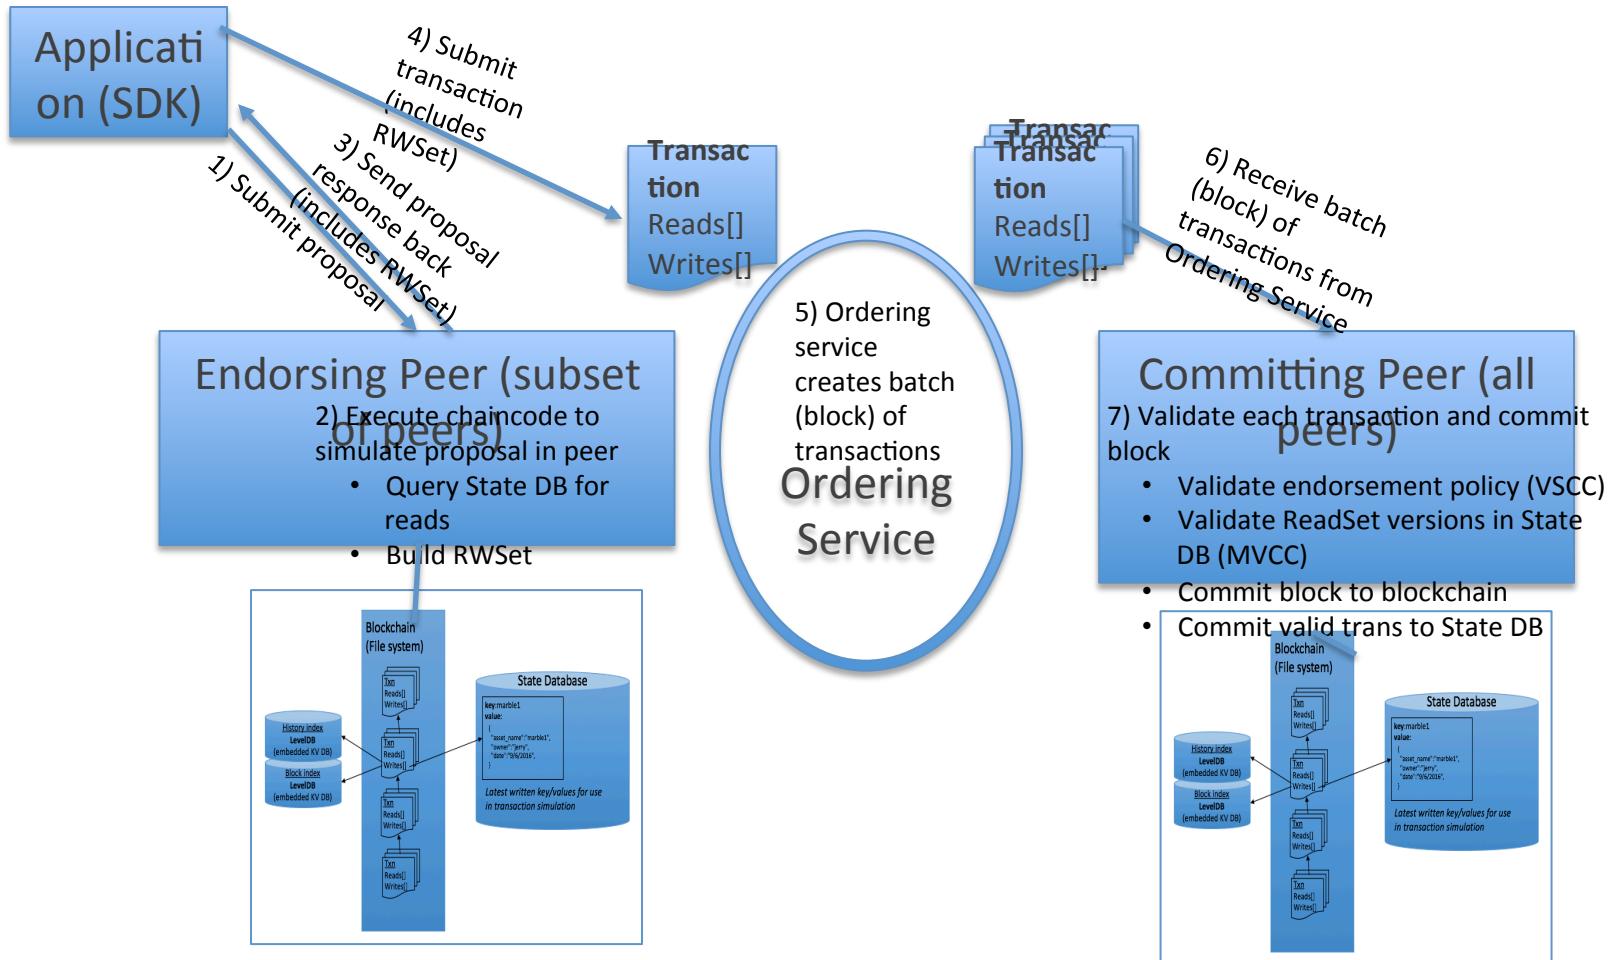
Immutable source of truth

'Index' of the blockchain for runtime queries

Ledger v1



v1 Transaction Lifecycle



Logical structure of a ReadWriteSet

```
Block{
    Transactions [
        {
            "Id" : txUUID2
            "Invoke" : "Method(arg1, arg2,..,argN)"
            "TxRWSet" : [
                { "Chaincode" : "ccId"
                    "Reads": [{"key" : "key1", "version" : "v1"}]
                    "Writes": [{"key" : "key1", "value" : bytes1}]
                } // end chaincode RWSet
            ] // end TxRWSet
        }, // end transaction with "Id" txUUID2

        { // another transaction },
    ] // end Transactions
} // end Block
```

Endorsing Peer (Simulation):

- Simulates transaction and generates ReadWriteSet

Committing Peer (Validation/Commit):

- Read set is utilized by MVCC validation check to ensure values read during simulation have not changed (ensures serializable isolation).
- Write set is applied to ledger (blockchain and state database) for validated transactions

Chaincode Data Patterns

Single key operations

GetState() / PutState() - Read and Write a single key/value.

Key range queries

Can be used in chaincode transaction logic (e.g. identify keys to update).

Fabric guarantees result set is stable between endorsement time and commit time, ensuring the integrity of the transaction.

GetStateByRange()

- Read keys between a startKey and endKey.

GetStateByPartialCompositeKey()

- Read keys that start with a common prefix. For example, for a chaincode key that is composed of K1-K2-K3 (composite key), ability to query on K1 or K1-K2 (performs range query under the covers). Replacement for v0.6 GetRows() table api.

Non-Key queries on data content beta in v1

Available when using a state database that supports content query (e.g. CouchDB)

Read-only queries against current state, not appropriate for use in chaincode transaction logic, unless application can guarantee result set is stable between endorsement time and commit time.

GetQueryResult()

- Pass a query string in the syntax of the state database

Pluggable state database - Objectives

Rich Query API for Blockchain (non-key query on data content)

- Leverage state-of-the-art database engines to extend query capabilities against blockchain data.
- Ensure interface supports plugging in different state database, for example by a vendor building on top of fabric
- To the degree possible, embed database and maintain within fabric, rather than requiring DBA skills

State Database options - Queryability

- In a **key/value database** such as **RocksDB** or **LevelDB**, the content is a blob and only queryable by key
 - Does not meet chaincode, auditing, reporting requirements for many use cases
- In a **document database** such as **CouchDB**, the content is JSON and fully queryable
 - Meets a large percentage of chaincode, auditing, and simple reporting requirements
 - For deeper reporting and analytics, replicate to an analytics engine such as Spark
 - Id/document data model compatible with existing chaincode key/value programming model, therefore no application changes are required when modeling chaincode data as JSON
- **SQL data stores** also possible, but requires more complicated relational transformation layer, as well as schema management.

State Database options - Queryability

- In a **key/value database** such as **RocksDB** or **LevelDB**, the content is a blob and only queryable by key
 - Does not meet chaincode, auditing, reporting requirements for many use cases
- In a **document database** such as **CouchDB**, the content is JSON and fully queryable
 - Meets a large percentage of chaincode, auditing, and simple reporting requirements
 - For deeper reporting and analytics, replicate to an analytics engine such as Spark
 - Id/document data model compatible with existing chaincode key/value programming model, therefore no application changes are required when modeling chaincode data as JSON
- **SQL data stores** also possible, but requires more complicated relational transformation layer, as well as schema management.

→ **Default in v1**

→ **Beta in v1**

→ **Future**

CouchDB State Database Details

- Single CouchDB instance under each peer
 - Redundancy provided by peer ‘replicas’ rather than database replicas
 - Normal blockchain model – peer with dedicated local data store
 - Since no database replicas, writes to database are guaranteed consistent and durable (not ‘eventually consistent’, as would be the case if there were database replicas)
 - Same basic model as when using local LevelDB state database
 - Except that CouchDB runs in separate local process, rather than embedded in peer process
- Configure for only local peer connections to CouchDB
 - Disable remote connections
 - Committing peer is only process that writes to state database
 - All queries go through peer authentication and authorization
 - Again, same interaction patterns as when using local LevelDB state database, but with more powerful query capability
 - CouchDB will be made available as Docker image. Docker compose will configure communication between peer container and couchdb containers (*coming Feb 2017*)

Pluggable State Database - The Challenges

How to support v1 endorsement/simulation model, when most databases do not support simulation result sets?

- That is, how to make uncommitted updates against an arbitrary database and determine the ReadWriteSet that is required for endorsement and commit validation? Not possible with most databases...

Solution:

- Query database for key values during endorser simulation, using database's rich query language
- Perform simulation updates in private workspace (peer memory) using normal chaincode APIs, e.g. PutState()
- Get ReadWriteSet from endorser simulation (Reads come from DB queries, Writes come from simulation in private workspace)
- Endorsement, Consensus, Validation use transaction ReadWriteSet Simulation Results as normal
- Apply Writes to database during Commit phase
- Peer maintains data integrity across blockchain file storage and state database

Transaction simulation is a proposal only. Updates are not yet applied to database. Implications:

- Transaction simulation does not support Read Your own Writes. GetState() always retrieves from state database.
- At commit time, validation is required to ensure conditions at simulation time (ReadSet) are still valid.
 - For key-based queries (GetState), validation step does a simple MVCC check on the ReadSet versions
 - For key-based range queries (GetStateByRange, GetStateByPartialCompositeKey) validation step re-queries to ensure result set is same – e.g. no added (phantom) items since simulation time
 - Non-key queries – Read-only queries against current state, not appropriate for use in chaincode transaction logic, unless application can guarantee result set is stable between endorsement time and commit time.

GetQueryResult() API

- Enabled when using a queryable state database such as CouchDB
 - Accepts a query string that gets passed to CouchDB state database `/_find` API, returns a set of key/values (documents)
 - Utilizes CouchDB `/_find` API query syntax
- Query API available in application chaincode
 - Read-only queries - *e.g. find all assets owned by Alice*
 - Client SDK invokes chaincode on an endorsing peer
 - Endorser returns response with ReadSet
 - Client does NOT submit transaction to ordering service (technically they could submit to ordering service to log the read, but this is not typical)
 - Queries as part of write transactions - *e.g. find all assets owned by Alice and transfer them to Bob*
 - Client SDK invokes chaincode on N endorsing peers
 - Endorser(s) return response with ReadWriteSet
 - Client SDK submits endorsed transaction to ordering service
 - Application must guarantee that result set will be stable between transaction simulation and commit time. If application cannot guarantee, structure data to use range queries or partial composite key queries instead
 - Access control enforced at either application or chaincode level

GetQueryResult() API- Indexes

- Indexes can be created in CouchDB to accelerate queries
- Initially, create indexes by calling CouchDB APIs from peer machine
- Evaluating options to create indexes upon chaincode deployment or via peer APIs

Query System Chaincode (QSCC)

- New system chaincode deployed by default in v1 to query blockchain
- Client can invoke against any peer, using same endorser request/response model that is used for application chaincode calls
- QSCC includes the following APIs (chaincode functions):
 - GetChainInfo
 - GetBlockByNumber
 - GetBlockByHash
 - GetTransactionByID

GetHistoryForKey

- `GetHistoryForKey()` API uses LevelDB history index to return history of values (states) for a key.
- Used for simple lineage/provenance scenarios.
- Available in application chaincode (similar to `GetQueryResult`)

Ledger Source Files

- core/ledger/ledger_interface.go
- core/ledger/kvledger/kv_ledger.go
- core/ledger/kvledger/txmgmt/statedb/statedb.go
- core/ledger/kvledger/txmgmt/statedb/statecouchdb/statecouchdb.go
- core/ledger/kvledger/history/historydb/historydb.go
- core/ledger/kvledger/history/historydb/historyleveldb/historyleveldb_query_executer.go

ledger_interface.go 1

```
93 // QueryExecutor executes the queries
94 // Get* methods are for supporting KV-based data model. ExecuteQuery method is for supporting a rich datamodel and query support
95 //
96 // ExecuteQuery method in the case of a rich data model is expected to support queries on
97 // latest state, historical state and on the intersection of state and transactions
98 type QueryExecutor interface {
99     // GetState gets the value for given namespace and key. For a chaincode, the namespace corresponds to the chaincodeId
100    GetState(namespace string, key string) ([]byte, error)
101    // GetStateMultipleKeys gets the values for multiple keys in a single call
102    GetStateMultipleKeys(namespace string, keys []string) ([][]byte, error)
103    // GetStateRangeScanIterator returns an iterator that contains all the key-values between given key ranges.
104    // startKey is included in the results and endKey is excluded. An empty startKey refers to the first available key
105    // and an empty endKey refers to the last available key. For scanning all the keys, both the startKey and the endKey
106    // can be supplied as empty strings. However, a full scan should be used judiciously for performance reasons.
107    // The returned ResultsIterator contains results of type *KV which is defined in protos/ledger/queryresult.
108    GetStateRangeScanIterator(namespace string, startKey string, endKey string) (commonledger.ResultsIterator, error)
109    // ExecuteQuery executes the given query and returns an iterator that contains results of type specific to the underlying data store.
110    // Only used for state databases that support query
111    // For a chaincode, the namespace corresponds to the chaincodeId
112    // The returned ResultsIterator contains results of type *KV which is defined in protos/ledger/queryresult.
113    ExecuteQuery(namespace, query string) (commonledger.ResultsIterator, error)
114    // GetPrivateData gets the value of a private data item identified by a tuple <namespace, collection, key>
115    GetPrivateData(namespace, collection, key string) ([]byte, error)
116    // GetPrivateDataMultipleKeys gets the values for the multiple private data items in a single call
117    GetPrivateDataMultipleKeys(namespace, collection string, keys []string) ([][]byte, error)
118    // GetPrivateDataRangeScanIterator returns an iterator that contains all the key-values between given key ranges.
119    // startKey is included in the results and endKey is excluded. An empty startKey refers to the first available key
120    // and an empty endKey refers to the last available key. For scanning all the keys, both the startKey and the endKey
121    // can be supplied as empty strings. However, a full scan shouold be used judiciously for performance reasons.
122    // The returned ResultsIterator contains results of type *KV which is defined in protos/ledger/queryresult.
123    GetPrivateDataRangeScanIterator(namespace, collection, startKey, endKey string) (commonledger.ResultsIterator, error)
124    // ExecuteQuery executes the given query and returns an iterator that contains results of type specific to the underlying data store.
125    // Only used for state databases that support query
126    // For a chaincode, the namespace corresponds to the chaincodeId
127    // The returned ResultsIterator contains results of type *KV which is defined in protos/ledger/queryresult.
128    ExecuteQueryOnPrivateData(namespace, collection, query string) (commonledger.ResultsIterator, error)
129    // Done releases resources occupied by the QueryExecutor
130    Done()
131 }
```

ledger_interface.go 2

```
133 // HistoryQueryExecutor executes the history queries
134 type HistoryQueryExecutor interface {
135     // GetHistoryForKey retrieves the history of values for a key.
136     // The returned ResultsIterator contains results of type *KeyModification which is defined in protos/ledger/queryresult.
137     GetHistoryForKey(namespace string, key string) (commonledger.ResultsIterator, error)
138 }
139
140 // TxSimulator simulates a transaction on a consistent snapshot of the 'as recent state as possible'
141 // Set* methods are for supporting KV-based data model. ExecuteUpdate method is for supporting a rich datamodel and query support
142 type TxSimulator interface {
143     QueryExecutor
144     // SetState sets the given value for the given namespace and key. For a chaincode, the namespace corresponds to the chaincodeId
145     SetState(namespace string, key string, value []byte) error
146     // DeleteState deletes the given namespace and key
147     DeleteState(namespace string, key string) error
148     // SetMultipleKeys sets the values for multiple keys in a single call
149     SetStateMultipleKeys(namespace string, kvs map[string][]byte) error
150     // ExecuteUpdate for supporting rich data model (see comments on QueryExecutor above)
151     ExecuteUpdate(query string) error
152     // SetPrivateData sets the given value to a key in the private data state represented by the tuple <namespace, collection, key>
153     SetPrivateData(namespace, collection, key string, value []byte) error
154     // SetPrivateDataMultipleKeys sets the values for multiple keys in the private data space in a single call
155     SetPrivateDataMultipleKeys(namespace, collection string, kvs map[string][]byte) error
156     // DeletePrivateData deletes the given tuple <namespace, collection, key> from private data
157     DeletePrivateData(namespace, collection, key string) error
158     // GetTxSimulationResults encapsulates the results of the transaction simulation.
159     // This should contain enough detail for
160     // - The update in the state that would be caused if the transaction is to be committed
161     // - The environment in which the transaction is executed so as to be able to decide the validity of the environment
162     //   (at a later time on a different peer) during committing the transactions
163     // Different ledger implementation (or configurations of a single implementation) may want to represent the above two pieces
164     // of information in different way in order to support different data-models or optimize the information representations.
165     // Returned type 'TxSimulationResults' contains the simulation results for both the public data and the private data.
166     // The public data simulation results are expected to be used as in V1 while the private data simulation results are expected
167     // to be used by the gossip to disseminate this to the other endorsers (in phase-2 of sidedb)
168     GetTxSimulationResults() (*TxSimulationResults, error)
169 }
```

kv_ledger.go

```
177 // NewTxSimulator returns new `ledger.TxSimulator`  
178 func (l *kvLedger) NewTxSimulator(txid string) (ledger.TxSimulator, error) {  
179     return l.txtmgmt.NewTxSimulator(txid)  
180 }  
181  
182 // NewQueryExecutor gives handle to a query executor.  
183 // A client can obtain more than one 'QueryExecutor's for parallel execution.  
184 // Any synchronization should be performed at the implementation level if required  
185 func (l *kvLedger) NewQueryExecutor() (ledger.QueryExecutor, error) {  
186     return l.txtmgmt.NewQueryExecutor(util.GenerateUUID())  
187 }  
188  
189 // NewHistoryQueryExecutor gives handle to a history query executor.  
190 // A client can obtain more than one 'HistoryQueryExecutor's for parallel execution.  
191 // Any synchronization should be performed at the implementation level if required  
192 // Pass the ledger blockstore so that historical values can be looked up from the chain  
193 func (l *kvLedger) NewHistoryQueryExecutor() (ledger.HistoryQueryExecutor, error) {  
194     return l.historyDB.NewHistoryQueryExecutor(l.blockStore)  
195 }  
196 }
```

statedb.go

```
23 // VersionedDB lists methods that a db is supposed to implement
24 type VersionedDB interface {
25     // GetState gets the value for given namespace and key. For a chaincode, the namespace corresponds to the chaincodeId
26     GetState(namespace string, key string) (*VersionedValue, error)
27     // GetVersion gets the version for given namespace and key. For a chaincode, the namespace corresponds to the chaincodeId
28     GetVersion(namespace string, key string) (*version.Height, error)
29     // GetStateMultipleKeys gets the values for multiple keys in a single call
30     GetStateMultipleKeys(namespace string, keys []string) ([]*VersionedValue, error)
31     // GetStateRangeScanIterator returns an iterator that contains all the key-values between given key ranges.
32     // startKey is inclusive
33     // endKey is exclusive
34     // The returned ResultsIterator contains results of type *VersionedKV
35     GetStateRangeScanIterator(namespace string, startKey string, endKey string) (ResultsIterator, error)
36     // ExecuteQuery executes the given query and returns an iterator that contains results of type *VersionedKV.
37     ExecuteQuery(namespace string, query string) (ResultsIterator, error)
38     // ApplyUpdates applies the batch to the underlying db.
39     // height is the height of the highest transaction in the Batch that
40     // a state db implementation is expected to use as a save point
41     ApplyUpdates(batch *UpdateBatch, height *version.Height) error
42     // GetLatestSavePoint returns the height of the highest transaction upto which
43     // the state db is consistent
44     GetLatestSavePoint() (*version.Height, error)
45     // ValidateKey tests whether the key is supported by the db implementation.
46     // For instance, leveldb supports any bytes for the key while the couchdb supports only valid utf-8 string
47     ValidateKey(key string) error
48     // BytesKeySupported returns true if the implementation (underlying db) supports the any bytes to be used as key.
49     // For instance, leveldb supports any bytes for the key while the couchdb supports only valid utf-8 string
50     BytesKeySupported() bool
51     // Open opens the db
52     Open() error
53     // Close closes the db
54     Close()
55 }
```

statecouchdb.go 1

```
141 // GetState implements method in VersionedDB interface
142 func (vdb *VersionedDB) GetState(namespace string, key string) (*statedb.VersionedValue, error) {
143     logger.Debugf(format: "GetState(). ns=%s, key=%s", namespace, key)
144
145     compositeKey := constructCompositeKey(namespace, key)
146
147     couchDoc, _, err := vdb.db.ReadDoc(string(compositeKey))
148     if err != nil {
149         return nil, err
150     }
151     if couchDoc == nil {
152         return nil, nil
153     }
154
155     // remove the data wrapper and return the value and version
156     returnValue, returnVersion := removeDataWrapper(couchDoc.JSONValue, couchDoc.Attachments)
157
158     return &statedb.VersionedValue{Value: returnValue, Version: returnVersion}, nil
159 }
```

```
844 func constructCompositeKey(ns string, key string) []byte {
845     compositeKey := []byte(ns)
846     compositeKey = append(compositeKey, compositeKeySep...)
847     compositeKey = append(compositeKey, []byte(key)...)
848     return compositeKey
849 }
850
851 func splitCompositeKey(compositeKey []byte) (string, string) {
852     split := bytes.SplitN(compositeKey, compositeKeySep, n: 2)
853     return string(split[0]), string(split[1])
854 }
```

statecouchdb.go 2

```
285 // ExecuteQuery implements method in VersionedDB interface
286 func (vdb *VersionedDB) ExecuteQuery(namespace, query string) (statedb.ResultsIterator, error) {
287     // Get the querylimit from core.yaml
288     queryLimit := ledgerconfig.GetQueryLimit()
289
290     queryString, err := ApplyQueryWrapper(namespace, query, queryLimit, querySkip: 0)
291     if err != nil {
292         logger.Debugf( format: "Error calling ApplyQueryWrapper(): %s\n", err.Error())
293         return nil, err
294     }
295
296     queryResult, err := vdb.db.QueryDocuments(queryString)
297     if err != nil {
298         logger.Debugf( format: "Error calling QueryDocuments(): %s\n", err.Error())
299         return nil, err
300     }
301
302     logger.Debugf( format: "Exiting ExecuteQuery")
303     return newQueryScanner(*queryResult), nil
304
305 }
```

query_wrapper.go#ApplyQueryWrapper 1

```
/*
ApplyQueryWrapper parses the query string passed to CouchDB
the wrapper prepends the wrapper "data." to all fields specified in the query
All fields in the selector must have "data." prepended to the field names
Fields listed in fields key will have "data." prepended
Fields in the sort key will have "data." prepended

- The query will be scoped to the chaincodeid
- limit be added to the query and is based on config
- skip is defaulted to 0 and is currently not used, this is for future paging implementation
```

In the example a contextID of "marble" is assumed.

Example:

Source Query:

```
{"selector": {"owner": {"$eq": "tom"}},
"fields": ["owner", "asset_name", "color", "size"],
"sort": ["size", "color"]}
```

Result Wrapped Query:

```
{"selector": {"$and": [{"chaincodeid": "marble"}, {"data.owner": {"$eq": "tom"}]}],
"fields": ["data.owner", "data.asset_name", "data.color", "data.size", "_id", "version"],
"sort": ["data.size", "data.color"], "limit": 10, "skip": 0}
```

*/

query_wrapper.go#ApplyQueryWrapper 2

```
64 func ApplyQueryWrapper(namespace, queryString string, queryLimit, querySkip int) (string, error) {
65
66     //create a generic map for the query json
67     jsonQueryMap := make(map[string]interface{})
68
69     //unmarshal the selected json into the generic map
70     decoder := json.NewDecoder(bytes.NewReader([]byte(queryString)))
71     decoder.UseNumber()
72     err := decoder.Decode(&jsonQueryMap)
73     if err != nil {
74         return "", err
75     }
76
77     //traverse through the json query and wrap any field names
78     processAndWrapQuery(jsonQueryMap)
79
80     //if "fields" are specified in the query, then add the "_id", "version" and "chaincodeid" fields
81     if jsonValue, ok := jsonQueryMap[jsonQueryFields]; ok {
82         //check to see if this is an interface map
83         if reflect.TypeOf(jsonValue).String() == "[]interface {}" {
84
85             //Add the "_id", "version" and "chaincodeid" fields, these are needed by default
86             jsonQueryMap[jsonQueryFields] = append(jsonValue.([]interface{}),
87             elems: "_id", "version", "chaincodeid")
88         }
89     }
90
91     //Check to see if the "selector" is specified in the query
92     if jsonValue, ok := jsonQueryMap[jsonQuerySelector]; ok {
93         //if the "selector" is found, then add the "$and" clause and the namespace filter
94         setNamespaceInSelector(namespace, jsonValue, jsonQueryMap)
95     } else {
96         //if the "selector" is not found, then add a default namespace filter
97         setDefaultNamespaceInSelector(namespace, jsonQueryMap)
98     }
99
100    //Add limit
101    jsonQueryMap[jsonQueryLimit] = queryLimit
102
103    //Add skip
104    jsonQueryMap[jsonQuerySkip] = querySkip
105
106    //Marshal the updated json query
107    editedQuery, _ := json.Marshal(jsonQueryMap)
108
109    logger.Debugf( format: "Rewritten query with data wrapper: %s", editedQuery)
110
111    return string(editedQuery), nil
112
113 }
```

statecouchdb.go#ApplyUpdates 1

```
307 // ApplyUpdates implements method in VersionedDB interface
308 func (vdb *VersionedDB) ApplyUpdates(batch *statedb.UpdateBatch, height *version.Height) error {
309
310     // STEP 1: GATHER DOCUMENT REVISION NUMBERS REQUIRED FOR THE COUCHDB BULK UPDATE
311     //           ALSO BUILD PROCESS BATCHES OF UPDATE DOCUMENTS BASED ON THE MAX BATCH SIZE
312
313     // initialize a missing key list
314     var missingKeys []*statedb.CompositeKey
315
316     //initialize a processBatch for updating bulk documents
317     processBatch := statedb.NewUpdateBatch()
318
319     //get the max size of the batch from core.yaml
320     maxBatchSize := ledgerconfig.GetMaxBatchUpdateSize()
321
322     //initialize a counter to track the batch size
323     batchSizeCounter := 0
324
325     // Iterate through the batch passed in and create process batches
326     // using the max batch size
327     namespaces := batch.GetUpdatedNamespaces()
328     for _, ns := range namespaces {
329         nsUpdates := batch.GetUpdates(ns)
330         for k, vv := range nsUpdates {
331
332             //increment the batch size counter
333             batchSizeCounter++
334
335             compositeKey := statedb.CompositeKey{Namespace: ns, Key: k}
336
337             // Revision numbers are needed for couchdb updates.
338             // vdb.committedDataCache.revisionNumbers is a cache of revision numbers based on ID
339             // Document IDs and revision numbers will already be in the cache for read/writes,
340             // but will be missing in the case of blind writes.
341             // If the key is missing in the cache, then add the key to missingKeys
342             _, keyFound := vdb.committedDataCache.revisionNumbers[compositeKey]
343
344             if !keyFound {
345                 // Add the key to the missing key list
346                 // As there can be no duplicates in UpdateBatch, no need check for duplicates.
347                 missingKeys = append(missingKeys, &compositeKey)
348             }
349
350             //add the record to the process batch
351             if vv.Value == nil {
352                 processBatch.Delete(ns, k, vv.Version)
353             } else {
354                 processBatch.Put(ns, k, vv.Value, vv.Version)
355             }
356 }
```

statecouchdb.go#ApplyUpdates 2

```
355 }  
356  
357 //Check to see if the process batch exceeds the max batch size  
358 if batchSizeCounter >= maxBatchSize {  
359  
360     //STEP 2: PROCESS EACH BATCH OF UPDATE DOCUMENTS  
361  
362     err := vdb.processUpdateBatch(processBatch, missingKeys) ←  
363     if err != nil {  
364         return err  
365     }  
366  
367     //reset the batch size counter  
368     batchSizeCounter = 0  
369  
370     //create a new process batch  
371     processBatch = statedb.NewUpdateBatch()  
372  
373     // reset the missing key list  
374     missingKeys = []*statedb.CompositeKey{}  
375  
376 }  
377  
378 }  
379 }  
380  
381 //STEP 3: PROCESS ANY REMAINING DOCUMENTS  
382 err := vdb.processUpdateBatch(processBatch, missingKeys)  
383 if err != nil {  
384     return err  
385 }  
386  
387 // STEP 4: IF THERE WAS SUCCESS UPDATING COUCHDB, THEN RECORD A SAVEPOINT FOR THIS BLOCK HEIGHT  
388  
389 // Record a savepoint at a given height  
390 err = vdb.recordSavepoint(height)  
391 if err != nil {  
392     logger.Errorf("Error during recordSavepoint: %s\n", err.Error())  
393     return err  
394 }  
395  
396 return nil  
397 }
```

statecouchdb.go#processUpdateBatch 1

```
399 //ProcessUpdateBatch updates a batch
400 func (vdb *VersionedDB) processUpdateBatch(updateBatch *statedb.UpdateBatch, missingKeys []*statedb.CompositeKey) error {
401
402     // An array of missing keys is passed in to the batch processing
403     // A bulk read will then add the missing revisions to the cache
404     if len(missingKeys) > 0 {
405
406         if logger.IsEnabledFor(logging.DEBUG) {
407             logger.Debugf("Retrieving keys with unknown revision numbers, keys= %s", printCompositeKeys(missingKeys))
408         }
409
410         err := vdb.LoadCommittedVersions(missingKeys)
411         if err != nil {
412             return err
413         }
414     }
415
416     // STEP 1: CREATE COUCHDB DOCS FROM UPDATE SET THEN DO A BULK UPDATE IN COUCHDB
417
418     // Use the batchUpdateMap for tracking couchdb updates by ID
419     // this will be used in case there are retries required
420     batchUpdateMap := make(map[string]*BatchableDocument)
421
422     namespaces := updateBatch.GetUpdatedNamespaces()
423     for _, ns := range namespaces {
424         nsUpdates := updateBatch.GetUpdates(ns)
425         for k, vv := range nsUpdates {
426             compositeKey := constructCompositeKey(ns, k)
427
428             // Create a document structure
429             couchDoc := &couchdb.CouchDoc{}
430
431             // retrieve the couchdb revision from the cache
432             // Documents that do not exist in couchdb will not have revision numbers and will
433             // exist in the cache with a revision value of nil
434             revision := vdb.committedDataCache.revisionNumbers[statedb.CompositeKey{Namespace: ns, Key: k}]
435
436             var isDelete bool // initialized to false
437             if vv.Value == nil {
438                 isDelete = true
439             }
440
441             logger.Debugf("Channel [%s]: key(string)=[%s] key(bytes)=[%#v], prior revision=[%s], isDelete=[%t]",
442                         vdb.dbName, string(compositeKey), compositeKey, revision, isDelete)
443
444             if isDelete {
445                 // this is a deleted record. Set the _deleted property to true
446                 couchDoc.JSONValue = createCouchdbDocJSON(string(compositeKey), revision, value: nil, ns, vv.Version, deleted: true)
447             } else {
448             }
```

statecouchdb.go#processUpdateBatch 2

```
448
449
450 } else {
451     if couchdb.IsJSON(string(vv.Value)) {
452         // Handle as json
453         couchDoc.JSONValue = createCouchdbDocJSON(string(compositeKey), revision, vv.Value, ns, vv.Version, deleted: false)
454     } else { // if value is not json, handle as a couchdb attachment
455
456         attachment := &couchdb.AttachmentInfo{}
457         attachment.AttachmentBytes = vv.Value
458         attachment.ContentType = "application/octet-stream"
459         attachment.Name = binaryWrapper
460         attachments := append([]*couchdb.AttachmentInfo{}, attachment)
461
462         couchDoc.Attachments = attachments
463         couchDoc.JSONValue = createCouchdbDocJSON(string(compositeKey), revision, value: nil, ns, vv.Version, deleted: false)
464
465     }
466
467
468     // Add the current document, revision and delete flag to the update map
469     batchUpdateMap[string(compositeKey)] = &BatchableDocument{CouchDoc: *couchDoc, Deleted: isDelete}
470
471 }
472
473
474 if len(batchUpdateMap) > 0 {
475
476     //Add the documents to the batch update array
477     batchUpdateDocs := []*couchdb.CouchDoc{}
478     for _, updateDocument := range batchUpdateMap {
479         batchUpdateDocument := updateDocument
480         batchUpdateDocs = append(batchUpdateDocs, &batchUpdateDocument.CouchDoc)
481     }
482
483     // Do the bulk update into couchdb
484     // Note that this will do retries if the entire bulk update fails or times out
485     batchUpdateResp, err := vdb.db.BatchUpdateDocuments(batchUpdateDocs)
486     if err != nil {
487         return err
488     }
489
490     // STEP 2: IF INDIVIDUAL DOCUMENTS IN THE BULK UPDATE DID NOT SUCCEED, TRY THEM INDIVIDUALLY
491
492     // iterate through the response from CouchDB by document
493     for _, respDoc := range batchUpdateResp {
494
495         // If the document returned an error, retry the individual document
496         if respDoc.Ok != true {
497
498             batchUpdateDocument := batchUpdateMap[respDoc.ID]
499
500             var err error
```

statecouchdb.go#processUpdateBatch 3

```
500
501
502     var err error
503
504     //Remove the "_rev" from the JSON before saving
505     //this will allow the CouchDB retry logic to retry revisions without encountering
506     //a mismatch between the "If-Match" and the "_rev" tag in the JSON
507     if batchUpdateDocument.CouchDoc.JSONValue != nil {
508         err = removeJSONRevision( jsonValue: &batchUpdateDocument.CouchDoc.JSONValue)
509         if err != nil {
510             return err
511         }
512     }
513
514     // Check to see if the document was added to the batch as a delete type document
515     if batchUpdateDocument.Deleted {
516
517         //Log the warning message that a retry is being attempted for batch delete issue
518         logger.Warningf( format: "CouchDB batch document delete encountered an problem. Retrying delete for document ID:%s", respDoc.ID)
519
520         // If this is a deleted document, then retry the delete
521         // If the delete fails due to a document not being found (404 error),
522         // the document has already been deleted and the DeleteDoc will not return an error
523         err = vdb.db.DeleteDoc(respDoc.ID, rev: "")
524     } else {
525
526         //Log the warning message that a retry is being attempted for batch update issue
527         logger.Warningf( format: "CouchDB batch document update encountered an problem. Retrying update for document ID:%s", respDoc.ID)
528
529         // Save the individual document to couchdb
530         // Note that this will do retries as needed
531         _, err = vdb.db.SaveDoc(respDoc.ID, rev: "", &batchUpdateDocument.CouchDoc)
532
533     }
534
535     // If the single document update or delete returns an error, then throw the error
536     if err != nil {
537
538         errorString := fmt.Sprintf( format: "Error occurred while saving document ID = %v  Error: %s  Reason: %s\n",
539             respDoc.ID, respDoc.Error, respDoc.Reason)
540
541         logger.Errorf(errorString)
542         return fmt.Errorf(errorString)
543     }
544
545     }
546
547     return nil
548 }
```

historydb.go

```
26 // HistoryDBProvider provides an instance of a history DB
27 type HistoryDBProvider interface {
28     // GetDBHandle returns a handle to a HistoryDB
29     GetDBHandle(id string) (HistoryDB, error)
30     // Close closes all the HistoryDB instances and releases any resources held by HistoryDBProvider
31     Close()
32 }
33
34 // HistoryDB – an interface that a history database should implement
35 type HistoryDB interface {
36     NewHistoryQueryExecutor(blockStore blkstorage.BlockStore) (ledger.HistoryQueryExecutor, error)
37     Commit(block *common.Block) error
38     GetLastSavepoint() (*version.Height, error)
39     ShouldRecover(lastAvailableBlock uint64) (bool, uint64, error)
40     CommitLostBlock(blockAndPvtdata *ledger.BlockAndPvtData) error
41 }
42 }
```

historyleveldb_query_executer.go 1

```
40 // GetHistoryForKey implements method in interface `ledger.HistoryQueryExecutor`
41 func (q *LevelHistoryDBQueryExecutor) GetHistoryForKey(namespace string, key string) (commonledger.ResultsIterator, error) {
42
43     if ledgerconfig.IsHistoryDBEnabled() == false {
44         return nil, errors.New(text: "History tracking not enabled - historyDatabase is false")
45     }
46
47     var compositeStartKey []byte
48     var compositeEndKey []byte
49     compositeStartKey = historydb.ConstructPartialCompositeHistoryKey(namespace, key, endkey: false)
50     compositeEndKey = historydb.ConstructPartialCompositeHistoryKey(namespace, key, endkey: true)
51
52     // range scan to find any history records starting with namespace~key
53     dbItr := q.historyDB.db.GetIterator(compositeStartKey, compositeEndKey)
54     return newHistoryScanner(compositeStartKey, namespace, key, dbItr, q.blockStore), nil
55 }
56
57 //historyScanner implements ResultsIterator for iterating through history results
58 type historyScanner struct {
59     compositePartialKey []byte //compositePartialKey includes namespace~key
60     namespace          string
61     key                string
62     dbItr              iterator.Iterator
63     blockStore         blkstorage.BlockStore
64 }
65
66 func newHistoryScanner(compositePartialKey []byte, namespace string, key string,
67 dbItr iterator.Iterator, blockStore blkstorage.BlockStore) *historyScanner {
68     return &historyScanner{compositePartialKey: compositePartialKey, namespace: namespace, key: key, dbltr: dbItr, blockStore: blockStore}
69 }
```

historyleveldb_query_executer.go 2

```
historyleveldb_query_executer.go x kv_ledger.go x config.go x couchdbutil.go x historydb.go x historyleveldb.go x util.go x val
71 func (scanner *historyScanner) Next() (commonledger.QueryResult, error) {
72     if !scanner.dbItr.Next() {
73         return nil, nil
74     }
75     historyKey := scanner.dbItr.Key() // history key is in the form namespace~key~blocknum~tranum
76
77     // SplitCompositeKey(namespace~key~blocknum~tranum, namespace~key~) will return the blocknum~tranum in second position
78     _, blockNumTranNumBytes := historydb.SplitCompositeHistoryKey(historyKey, scanner.compositePartialKey)
79     blockNum, bytesConsumed := util.DecodeOrderPreservingVarUint64(blockNumTranNumBytes[0:])
80     tranNum, _ := util.DecodeOrderPreservingVarUint64(blockNumTranNumBytes[bytesConsumed:])
81     logger.Debugf(format: "Found history record for namespace:%s key:%s at blockNumTranNum %v:%v\n",
82         scanner.namespace, scanner.key, blockNum, tranNum)
83
84     // Get the transaction from block storage that is associated with this history record
85     tranEnvelope, err := scanner.blockStore.RetrieveTxByBlockNumTranNum(blockNum, tranNum)
86     if err != nil {
87         return nil, err
88     }
89
90     // Get the txid, key write value, timestamp, and delete indicator associated with this transaction
91     queryResult, err := getKeyModificationFromTran(tranEnvelope, scanner.namespace, scanner.key)
92     if err != nil {
93         return nil, err
94     }
95     logger.Debugf(format: "Found historic key value for namespace:%s key:%s from transaction %s\n",
96         scanner.namespace, scanner.key, queryResult.(*queryresult.KeyModification).TxId)
97     return queryResult, nil
98 }
99
100 func (scanner *historyScanner) Close() {
101     scanner.dbItr.Release()
102 }
```

historyleveldb_query_executer.go 3

```
historyleveldb_query_executer.go x kv_ledger.go x config.go x couchdbutil.go x historydb.go x historyleveldb.go x util.go x validator.go x
104 // getTxIDandKeyWriteValueFromTran inspects a transaction for writes to a given key
105 func getKeyModificationFromTran(tranEnvelope *common.Envelope, namespace string, key string) (commonledger.QueryResult, error) {
106     logger.Debugf( format: "Entering getKeyModificationFromTran()\n", namespace, key)
107
108     // extract action from the envelope
109     payload, err := putils.GetPayload(tranEnvelope)
110     if err != nil {
111         return nil, err
112     }
113
114     tx, err := putils.GetTransaction(payload.Data)
115     if err != nil {
116         return nil, err
117     }
118
119     _, respPayload, err := putils.GetPayloads(tx.Actions[0])
120     if err != nil {
121         return nil, err
122     }
123
124     chdr, err := putils.UnmarshalChannelHeader(payload.Header.ChannelHeader)
125     if err != nil {
126         return nil, err
127     }
128
129     txID := chdr.TxId
130     timestamp := chdr.Timestamp
131
132     txRWSet := &rwsetutil.TxRwSet{}
```

historyleveldb_query_executer.go 4

```
133
134     // Get the Result from the Action and then Unmarshal
135     // it into a TxReadWriteSet using custom unmarshalling
136     if err = txRWSet.FromProtoBytes(respPayload.Results); err != nil {
137         return nil, err
138     }
139
140     // look for the namespace and key by looping through the transaction's ReadWriteSets
141     for _, nsRWSet := range txRWSet.NsRwSets {
142         if nsRWSet.NameSpace == namespace {
143             // got the correct namespace, now find the key write
144             for _, kvWrite := range nsRWSet.KvRwSet.Writes {
145                 if kvWrite.Key == key {
146                     return &queryresult.KeyModification{TxId: txID, Value: kvWrite.Value,
147                                         Timestamp: timestamp, IsDelete: kvWrite.IsDelete}, nil
148                 }
149             } // end keys loop
150             return nil, errors.New( text: "Key not found in namespace's writeset")
151         } // end if
152     } //end namespaces loop
153     return nil, errors.New( text: "Namespace not found in transaction's ReadWriteSets")
154
155
156 }
```

Final Thoughts + Q&A

- I hope you have enjoyed today's talk and learned how a Hyperledger Peer endorses, validates, simulates and queries transactions.